



Behaviour-based detection of Visual Interaction Obstacles with 1D and 2D Convolutional Neural Networks

Bachelor Thesis at the Cognitive Systems Lab
Prof. Dr.-Ing. Tanja Schultz
Faculty 3: Mathematics and Computer Science
University of Bremen

by

Anthony Mendil

Supervisor:

Mazen Salouz

Examiners:

Dr. Felix Putze

Prof. Dr. Sebastian Siebertz

Day of registration: August 24, 2020

Day of submission: September 29, 2020

I hereby declare that I am writing this work independently and have not used any sources or resources other than those specified.

Bremen, the September 29, 2020

Zusammenfassung

Die Leistungsfähigkeit von Menschen in der Interaktion mit Computern kann durch verschiedene Hindernisse beeinträchtigt werden. Man bezeichnet sie als Hindernisse der Mensch-Computer-Interaktion. Es ist möglich kognitive adaptive Systeme zu implementieren, die diese Hindernisse durch geeignete Anpassungen der Benutzerschnittstelle kompensieren. Damit eine solche Kompensation möglich ist, müssen solche Hindernisse jedoch zuerst erkannt werden. In dieser Arbeit wird ein flüchtiges visuelles Interaktionshindernis namens Blendeffekt auf der Grundlage von Verhaltensdaten von 20 Testpersonen erkannt, die ein Memory Spiel spielen, bei dem das Ziel ist Karten mit gleichen Farben zu finden. Der Blendeffekt beschreibt ein Szenario, in dem Sonnenlicht auf den Bildschirm scheint und somit zu einer geringeren Unterscheidbarkeit der Farben führt. Zur Erkennung dieses Hindernisses werden ein- und zweidimensionale convolutional neural networks (1D convnets und 2D convnets) verwendet. Während 1D convnets ein üblicher Ansatz für Sequenz-Klassifikationen sind, werden 2D convnets eher im Bereich der Bildverarbeitung eingesetzt. In dieser Arbeit wird jedoch ein neuartiger Ansatz implementiert, der darin besteht, aus den sequentiellen Daten synthetische Bilder zu erzeugen und 2D convnets zur Klassifikation dieser Bilder zu verwenden. Entsprechend wird ein ursprüngliches Sequenz-Klassifikationsproblem in den Bereich der Bildverarbeitung überführt. Um die Menge verfügbarer Trainingsdaten zu erhöhen, wird ein kognitiver Benutzersimulator verwendet, der unter Verwendung eines generativen Optimierungsalgorithmus neue Verhaltensdaten simuliert. Die Performanz der 1D und 2D convnets wird im Hinblick auf die Menge der simulierten Daten und die Anzahl der im Training enthaltenen Spielzüge analysiert. Darüber hinaus werden Änderungen am Simulator vorgenommen und die Performanz der Modelle vor und nach diesen Änderungen wird verglichen. Die Ergebnisse zeigen, dass 1D convnets bei der Erkennung des Hindernisses nach 5 und 10 Spielzügen bessere Ergebnisse erzielen als 2D convnets, während das Gegenteil nach 15 und 20 Zügen der Fall ist. Es werden vier ensemble-basierte Systeme implementiert, jeweils eines für 5, 10, 15 und 20 Züge. Die ersten beiden basieren auf 1D und die beiden anderen auf 2D convnets. Jedes dieser Systeme besteht aus mehreren Modellen, die abstimmen um die endgültige Entscheidung zu treffen. Die Genauigkeit dieser Systeme in der Reihenfolge der Anzahl der Züge beträgt 72,5%, 82,5%, 80% und 85%.

Abstract

The performance of humans interacting with computers can be impaired by several obstacles. Such obstacles are called Human Computer Interaction (HCI) obstacles. Cognitive adaptive systems can be implemented that compensate those obstacles with suitable User Interface (UI) adaptations. However, such obstacles first need to be detected. In this work, a transient visual HCI interaction obstacle called glare effect is detected based on behavioural data of 20 subjects playing a matching pairs game where the cards have different colours. The glare effect describes a scenario in which sunlight shines onto the display, resulting in less distinguishability of the colours. For the detection of this obstacle one and two dimensional convolutional neural networks (1D convnets and 2D convnets) are utilized. While 1D convnets are a common approach for sequence classifications, 2D convnets are generally more used in the area of computer vision. However, in this work a novel approach is implemented that consists of creating synthetic images out of the sequential data and using 2D convnets to classify the images, thus turning what was originally a sequence classification task into an image classification task. In order to increase the available training data a cognitive user simulator is used that implements a generative optimization algorithm to simulate behavioural data. The performance of the 1D and 2D convnets is analysed in regard to the the amount of simulated data and the number of game rounds included in the training. Furthermore, changes to the simulator are made and the performance of the models before and after the changes are compared. Results show that 1D outperform 2D convnets in detecting the obstacle after 5 and 10 rounds of the game, while 2D outperform 1D convnets in detecting it after 15 and 20 rounds. Four ensemble-based systems are implemented, one each for 5, 10, 15 and 20 rounds. The first two are based on 1D and the other two on 2D convnets. Each system consists of multiple models that vote for the final prediction. The accuracies of these systems in the order of the number of rounds are 72.5%, 82.5%, 80% and 85%.

Contents

1	Introduction	1
2	Memory Game	3
3	CMM-based Cognitive User Simulation	5
4	Collected Data	7
5	Data Preparation	9
5.1	Similarity matrix creation	10
5.1.1	Conceptualization	10
5.1.2	Screenshot extraction	14
5.1.3	Implementation	15
5.1.4	Testing	16
5.1.5	Final Matrix	17
5.2	Incorporation of the new similarity matrix	18
5.3	Removal of invalid logs and balancing of the data set	18
5.4	Simulation of user behaviour	19
5.5	Sorting logs by quality	19
5.6	Evaluation of simulation	20
5.6.1	Results before and after changing the simulator	20
5.6.2	Using a subset of the data	24
5.6.3	Conclusion of the evaluation	27
5.7	Feature generation	28
5.7.1	1D convnet features	28
5.7.2	2D convnet features	29
6	Convolutional neural networks	33
6.1	Layers	34
6.1.1	Convolution layers	34
6.1.2	Max-Pooling layers	36
6.1.3	Dropout layers	36
6.1.4	Flatten layers	36
6.1.5	Dense Layers	36
6.2	Activation functions	37
6.3	Receptive field	37
6.4	Differences of one dimensional convolutional neural networks	38

7	Training neural networks	39
8	Classification	41
8.1	Models	41
8.1.1	1D convnet	42
8.1.2	2D convnet	43
8.2	Training setup	44
8.3	Training results and evaluation	45
8.4	Ensemble-based systems	50
9	Conclusion	53
9.1	Prospect	54
A	Appendix	55
A.1	Scripts	55
A.2	Data	57
	Bibliography	59

List of Figures

2.1	Screenshot of the game without any obstacles.	4
2.2	Screenshot of the game field with simulated sunlight.	4
5.1	Screenshot of the game field with simulated sunlight.	10
5.2	Exemplary showing the number of different comparisons of two colours.	11
5.3	Final similarity matrix for the glare effect.	17
5.4	Comparison of matching pairs per round of simulated and real glare effect an no obstacle games. Before the changes to the simulator.	20
5.5	Comparison of penalties per round of simulated and real glare effect an no obstacle games. Before the changes to the simulator.	20
5.6	Comparison of matching pairs per round of simulated and real glare effect an no obstacle games. After the changes to the simulator.	21
5.7	Comparison of penalties per round of simulated and real glare effect an no obstacle games. After the changes to the simulator.	21
5.8	Comparison of matching pairs per round of simulated and real glare effect and no obstacle games in the first 10 rounds. Only the best 15 simulations per real game were used. Furthermore, the data used is from after the changes to the simulator.	26
5.9	Comparison of penalties per round of simulated and real glare effect and no obstacle games in the first 10 rounds. Only the best 15 simulations per real game were used. Furthermore, the data used is from after the changes to the simulator.	26
5.10	Synthetic image created using the statistical features. Origin of the image in the top left corner.	30
5.11	Synthetic image created using the statistical features. Origin of the image in the bottom left corner.	30
5.12	Card code section of a synthetic image for a no obstacle game.	32
5.13	Card code section of a synthetic image for a glare effect game.	32
6.1	Valid locations of 3 x 3 windows in a feature map.	34
6.2	The whole process of a two dimensional convolution.	35
6.3	Visualization of a linearly growing receptive field in stacked convolutions.	37
6.4	The whole process of a one dimensional convolution.	38
7.1	Core concepts in training neural networks.	39
8.1	Visualization of the utilized 1D convnet topology.	42
8.2	Visualization of the utilized 2D convnet topology.	43
8.3	Best results for the different number of rounds before the changes to the simulator.	47

8.4	Average training and test accuracy of 400 1D convnets on 5 rounds and a ratio of 9 simulated games for each real game.	49
8.5	Average training and test accuracy of 400 1D convnets on 10 rounds using no simulated games.	49
8.6	Average training and test accuracy of 400 2D convnets on 15 rounds and a ratio of 9 simulated games for each real game.	49
8.7	Average training and test accuracy of 400 2D convnets on 20 rounds and a ratio of 4 simulated games for each real game.	49

List of Tables

4.1	Static assignment of colours to numbers	7
4.2	Static and dynamic mapping of card codes.	8
5.1	Results of paired sample t-tests for different comparisons of matching pairs per round before the changes to the simulator. All 1000 simulated games per real game were used.	23
5.2	Results of paired sample t-tests for different comparisons of penalties per round before the changes to the simulator. All 1000 simulated games per real game were used.	23
5.3	Results of paired sample t-tests for different comparisons of matching pairs per round after the changed to the simulator. All 1000 simulated games per real game were used.	23
5.4	Results of paired sample t-tests for different comparisons of penalties per round after the changes to the simulator. All 1000 simulated games per real game were used.	23
5.5	Results of paired sample t-tests for different comparisons of matching pairs per round before the changes to the simulator. Only the best 15 simulated logs from each real log and the first 10 rounds are used. . .	25
5.6	Results of paired sample t-tests for different comparisons of penalties per round before the changes to the simulator. Only the best 15 simulated logs from each real log and the first 10 rounds are used. . .	25
5.7	Results of paired sample t-test for different comparisons of matching pairs per round after the changes to the simulator. Only the best 15 simulated logs from each real log and the first 10 rounds are used. . .	25
5.8	Results of paired sample t-test for different comparisons of penalties per round after the changes to the simulator. Only the best 15 simulated logs from each real log and the first 10 rounds are used. . .	25
5.9	Mean values of matching pairs for real and simulated glare effect games for the first 10 rounds. Only the 15 best simulations per real game are used.	27
5.10	Range and vertical space of the different component in the synthetic image.	31
8.1	Best accuracies and losses produced by the 1D convnet, prior to the changes to the simulator.	46
8.2	Best accuracies and losses produced by the 2D convnet, prior to the changes to the simulator.	46
8.3	Best accuracies and losses produced by the 1D convnet, after the changes to the simulator.	46

8.4	Best accuracies and losses produced by the 2D convnet, after the changes to the simulator.	46
8.5	Highest accuracies before and after the changes to the simulator. . . .	48
8.6	The chosen configurations and their accuracies.	48
8.7	Results of voting based systems compared to average of the models. .	51

1. Introduction

More and more people use electronic devices in their lives. Alone the number of smartphone users in Germany was close to ten times higher in 2019 compared to 2009 [Ten20]. As a result it becomes increasingly important to ensure a good interaction experience between humans and electronic devices. A good interaction experience with a device is not only characterised by its ease of use, intuitiveness, robustness and security, but also that it can be used by everybody and in case of mobile devices at every place. Who can use the device and where it can be used is influenced by so-called interaction obstacles. Two types of such interaction obstacles are visual and memory based obstacles. While a memory obstacle could for instance be caused by a secondary task workload, visual obstacles can for example be caused by a red-green colour vision deficiency [SPIS19, p. 1]. If a human computer interaction is based on differentiating colours, the interaction experience and performance of users with a red-green colour vision deficiency will likely be heavily harmed if no assistance is provided. The red-green colour vision deficiency is a persistent interaction obstacle that poses an issue during the whole interaction, whereas so-called transient interaction obstacles only occur temporarily.

In this work, a transient visual interaction obstacle called glare effect will be addressed. It describes a scenario in which sunlight shines onto the display, resulting in less distinguishability of the colours. Therefore, the user might not be capable to successfully interact with the device. Especially the interaction performance with mobile devices such as smartphones or tablets can be significantly impaired by the glare effect, as these devices are often used in the open. In order to compensate such obstacles cognitive adaptive systems can be implemented that first detect the obstacle and then provide a suitable user interface (UI) adaptation. Although not part of this work, an example UI adaptation for compensating the glare effect could be using shapes or sounds instead of colours as the differentiating factor.

Based on behavioural data collected using a matching pairs game, the aim of this work is to detect the visual interaction obstacle being the glare effect. The game can be played on android devices and was developed at the cognitive system lab (CSL). A game consists of 14 cards of which two cards each have the same colour and the goal is to find all 7 pairs. There are various different game modes, but the only

two necessary for this work are the mode without any obstacles and the glare effect mode which imitates the effect of sunlight shining onto the display. The data used had already been collected through an user study. As the number of participants was very limited, a cognitive user simulator, also implemented at the CSL, is used to simulate additional behavioural data. In order to successfully simulate the glare effect the simulator requires a so-called similarity matrix that describes the difference of the colours of the cards under the influence of the sunlight shining onto the display. Therefore, one step is to create such a similarity matrix. Once the simulation of user behaviour is completed the real and the simulated data are used to determine whether subjects played a game without any obstacles or a game affected by the glare effect. This decision is made using one and two dimensional convolutional neural networks (1D convnets and 2D convnets). The 1D convnets use statistical features for each step in the game (one round consists of flipping two cards), while the 2D convnets use synthetic images created with those statistical features. These models are trained and tested for different configurations regarding the ratio of simulated to real data and the number of game rounds included. The results of the different configurations and models are compared and the best performing ones are selected to form the bases of ensemble-based systems. In total four ensemble-based systems, consisting of many models voting for the final prediction, are implemented for detecting the glare effect after different numbers of game rounds.

2. Memory Game

The game used for the data collection is a matching pairs game written in Kotlin for Android devices. The original version was developed by Rafael Miranda in the context of a bachelor thesis [Mir18]. Originally, the game consisted of cards picturing animals and was only used to model general poor eyesight. Later it was expanded by adding a coloured card set to also model red-green colour vision deficiency [Ihr19, p. 7]. In the following the current functionalities of the game will be explained, only covering those related to the coloured cards. One game consists of 14 cards and each round consists of two cards being turned face up. If the cards match, they will be removed from the field. Otherwise they are turned face down again. The player wins once all 7 pairs have been discovered.

The game differentiates between two groups of obstacles: memory obstacles and visual obstacles. Additionally, there is a classic mode providing the possibility to play without any obstacles which uses the red-green card set, shown in figure 2.1. The memory obstacle mode consists of the same card set and an additional side task: Every time a card is flipped a random number between 0 and 10 is called out and the player must add all these numbers. After the final pair was found the sum of all numbers is requested. The game contains modes for two types of visual obstacles: colour blindness and the glare effect. Colour blindness, more precisely a red-green colour vision deficiency, is simulated by replacing the red-green card set with a brown shifted card set. The glare effect describes a scenario in which a light source shines onto the display, reflects from it and makes it more difficult to differentiate the colours of the cards. This effect is imitated by the game and can be seen in figure 2.2.

To help in case of interaction obstacles the game provides multiple ways of assistance. Although they are not relevant for this thesis as only data is used where subjects had no assistance, they will be explained briefly. There are two types of assistance: memory and visual assistance. Memory assistance is provided by flipping the cards from the previous turn once two non matching cards are flipped. After a short period of time they are turned face down again. If the game is played with visual assistance, each pair gets assigned a letter that is called out once a card is flipped. This introduces a new way of orientation that does not rely on the sense of sight, but

instead on the sense of hearing. There is a game mode for each combination of obstacle (no obstacle, memory obstacle, colour blindness, glare effect) and assistance (no assistance, memory assistance, visual assistance). There were also other assistances implemented such as increasing the size of cards if they are revealed, but they were only used for tests. The glare effect no assistance and the no obstacle no assistance modes are the only ones relevant in this thesis. Figure 2.1 and 2.2 show screenshots in the two modes.




Figure 2.1: Screenshot of the game without any obstacles. All cards are turned face up.



Figure 2.2: Screenshot of the game field with simulated sunlight. Sunbeams are especially visible on the right hand side. All cards are turned face up.

3. CMM-based Cognitive User Simulation

The cognitive system lab developed the computer program cognitive memory model (CMM). It is a general computational cognitive model of human memory inspired by the ACT-R theory [SP18a, p. 161]. ACT-R stands for adaptive control of thought-rational and is a cognitive architecture that aims to explain how the components of the mind work to produce coherent cognition [ABB⁺04, p. 1036]. In the context of the matching pairs game the CMM was used to implement a generative model. By modelling memorized and forgotten revealed cards and deciding in each turn whether to explore unknown cards or to exploit the gathered knowledge to reveal matching pairs, the course of a matching pairs game can be simulated taking the capability of human memory into account [SP18b, p. 161]. Furthermore, it is possible to define a similarity matrix to simulate similarities between cards in matching pairs games. This matrix includes similarity values between 0 and 1 for each colour combination of cards. Such a similarity matrix can be used in the context of visual interaction obstacles to emulate human confusion. In related research it was used to simulate the confusion caused by a red-green colour vision deficiency and the usage of a red-green card set [SPIS19, p. 3]. For this thesis the only two modes that are relevant are the glare effect and the no obstacle modes, both without any assistance. In order to simulate no obstacle games no similarity matrix is needed, but for glare effect games a new similarity matrix is needed and therefore created in section 5.1 [Similarity matrix creation](#).

To accurately simulate the performance of human memory multiple parameters are randomly initialized and then repeatedly optimized using a genetic optimization algorithm. These parameters can be further extended. The genetic algorithm works by repeatedly simulating game sessions and updating the parameters by applying mutation and selection operations. To select the best parameter values two performance measurements are defined and used to compare the simulated game sessions and a real game as reference. These performance measurements are the number of matching pairs and the number of penalties per round. Penalties are given if a card is revealed whose partner has already been seen before but the pair was not picked up. The optimization of the parameter population is repeated over many

generations so that the performance in the simulated games best fits that in the real game. For example, one of the parameters optimized by the genetic optimization algorithm is the *similarity decay*. It reduces the similarity effects during the course of the game, since the user may learn to adapt to the visual interaction obstacle and therefore be less effected by it. Furthermore, the similarity effect is continually reduced as there are less cards in the game after pairs were discovered [SPIS19, p. 5].

As input the simulator takes a number of game logs and simulated a specified amount of new sessions using the genetic optimization algorithm. For instance, 20 game logs can be passed to simulate 1000 new sessions out of each game log, resulting in 20000 simulated games. The data contained in the game logs is shown in chapter 4 [Collected Data](#). In the context of simulating no obstacle games changes to the simulator were made that are explained in section [5.6.1 Results before and after changing the simulator](#).

4. Collected Data

While the subjects played the memory game, behavioural data and brain activity data was collected. Relevant for this work is only the behavioural data recorded in glare effect and no obstacle games. The eeg data is not used, because the aim is to exploit behavioural patterns of natural HCI situations without using any additional sensors. Therefore, eeg data will not be further explained. The data was collected from 22 subjects. They were university students between the age of 18 and 28. Each subject except one recorded two no obstacle and one glare effect game. The one exception did not record any glare effect games. As a result, there are 44 no obstacle and 21 glare effect games. The data was already provided and was not collected during this bachelor thesis.

In order to record behavioural data, each of the 14 cards was initially given a fixed card code and it was recorded which pairs of cards were turned face up in each round. The card code consists of two numbers separated by a dot. The first number specifies the colour of the card (between 1 and 7 as there are 7 colours) and the second number specifies if it is the first or second card with that colour. Once a game is started, all cards are shuffled. The initial assignment of colours to numbers between 1 and 7 is shown in table 4.1.

Colour	Number
Dark green	1
Brown	2
Red	3
Light green	4
Green	5
Orange	6
Dark red	7

Table 4.1: Static assignment of colours to numbers

The behavioural data was saved in logs, consisting of the card codes of each round followed by the timestamps of when the cards were flipped. Data of 20 rounds and therefore at maximum 40 flipped cards was recorded. If the game was completed in less than 20 turns, the remaining card code entries were filled with 0.0 and the remaining timestamp entries were filled with 0. An example of the sequence of card codes recorded during a game can be seen in the top half of table 4.2. The timestamps are irrelevant for this work and therefore not shown.

With the assignments shown in table 4.1, the first card code in the upper sequence of table 4.2 means that the first card to be flipped was the second green card. In the same turn, the second card that was turned around was the second brown card. This mapping is static as colours have the same numbers across all recorded games. However, the simulator requires a dynamic mapping of the card codes. By dynamic mapping of the card codes is meant, that the colours are not assigned to fixed numbers but that the numbers are assigned in the reveal order specific to each game. This becomes clear, by looking the card code sequence from above, but dynamically mapped. This can be seen in the bottom half of table 4.2.

Static mapping: 5.2,2.2,4.2,4.1,6.1,7.2,6.2,6.1,1.2,1.1, 7.1,2.2,2.1,7.1,5.2,5.1,3.1,2.1,7.1,3.1, 7.2,7.1,3.2,2.2,2.1,2.2,3.1,3.2,0.0,0.0, 0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0
Dynamic mapping: 1.1,2.1,3.1,3.2,4.1,5.1,4.2,4.1,6.1,6.2, 5.2,2.1,2.2,5.2,1.1,1.2,7.1,2.2,5.2,7.1, 5.1,5.2,7.2,2.1,2.2,2.1,7.1,7.2,0.0,0.0, 0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0

Table 4.2: Static and dynamic mapping of card codes

In this dynamic mapping of the card codes, each entry likewise consists of two numbers that are separated by a dot. However, the numbers are differently interpreted. The first number specifies what number of colour it is to be revealed and the second number specifies if it was the first or the second card of that colour. In the example above the colour green is the first to be revealed and therefore assigned to the value 1. As it is the first green card revealed the second component of the card code is 1 as well. As a result, the first entry is 1.1. Then a brown card is turned face up, meaning the second entry is 2.1 since it was the first brown card. Once the other brown card is discovered, the entry for that step will be 2.2. This pattern continues throughout the whole sequence. This means that the mapping is specific to each game and depends on the reveal order of the cards. These card codes are the foundation for the features calculated in section 5.7 Feature generation.

Logs with remapped card codes were already provided, but an additional component in the glare effect logs is needed in order to successfully simulate new games. The logs are used in the simulator to generate new games. When simulating the glare effect, the simulator requires a similarity matrix, that that correctly describes the colour differences of the cards under the influence of the glare effect, which does not exist yet. The creation of this matrix is done in section 5.1 Similarity matrix creation and its values are added to the glare effect logs in section 5.2 Incorporation of the new similarity matrix. Contrary to the glare effect logs, in the no obstacle logs, no changes need to be made, as the simulator does not use any similarity values when simulating no obstacle games. It should also be noted that the provided logs with remapped card codes additionally include four statistical features for the last turn. These consist of the number of remaining cards, the number of never revealed cards, the highest number of times the same card was revealed and the number of rounds since all pairs were found. These four values are ignored, as they are newly calculated for every step as explained in section 5.7.1 1D convnet features. Last but not least all logs end with a label, describing in which game mode the log was created.

5. Data Preparation

To prepare the data for the training many steps were necessary. One big step is to increase the available data by simulating user behaviour. Therefore, the simulator explained in chapter 3 [CMM-based Cognitive User Simulation](#) is used. For simulating games with no obstacles there is no need for a similarity matrix. However, in order to successfully simulate glare effect games, a special similarity matrix is required. Therefore, the first step is to create a similarity matrix that describes the differences of colours under the influence of the glare effect. Once completed, it is added to the glare effect logs. As not all logs are valid and can be used in the simulator, the invalid ones must be removed. The valid logs are then used to simulate user behaviour and thereby create new logs. These simulated logs are sorted by their quality and the simulation is evaluated. The evaluation inspired some changes to the simulator, after which the games are simulated again. The quality of the simulations before and after the changes to the simulator are analysed. Lastly, to prepare the data for training, the features for the 1D convnet and the 2D convnet are created.

5.1 Similarity matrix creation

5.1.1 Conceptualization

The aim is to create a similarity matrix that describes the differences between colours under the influence of the simulated sun light. The difficulty hereby is that the intensity of the light is not spread evenly across the whole field. Instead, as it would be if a real sun would shine onto the display, a certain area has the highest intensity and the intensity is lower the further away from that area. Additionally, there are lines of higher intensity coming from the brightest area, that simulate sunbeams. This influence of the simulated sunlight can be seen in figure 5.1.



Figure 5.1: Screenshot of the game field with simulated sunlight. Sunbeams are especially visible on the right hand side. All cards are turned face up.

As a result, simply extracting the RGB values for one pixel of each card in one game and calculating the differences has two major problems: Firstly, the differences are highly influenced by the position of the specific cards. If the matrix is calculated using a single glare effect game the differences are only representative for exactly that game and may be completely different if the cards are positioned differently. Secondly, extracting single pixels for each card in an image and using them to calculate the similarity values does in some cases not accurately describe the perceived colour difference. Depending on the position of the card on the field, its colours can significantly vary on different locations on the card due to the influence of sunbeams. As a result, it is possible that the difference of two cards is perceivable to humans because parts of the cards are not as strongly influenced by the sunlight. Extracting only single pixels directly in the sunbeams and using them to calculate the colour difference would in that case not produce representative results. As this behaviour is undesired, a different approach than simply extracting single pixels out of one game was needed.

The problem of the positional influence can be solved by creating similarity matrices for more than one game and calculating the mean colour differences of those matrices. The idea is that by using enough games so that all or most combinations of positions and colours are included and calculating the mean of all comparisons for each combination of colours, the result will not be influenced by the position of the cards. However, first needs to be determined which number of games satisfies this condition. Each game consists of 14 cards, with each two cards having the same colours. When determining the difference between two colours there are two different cases that are also visualized in figure 5.2:

- 1. The colours are not the same: In each game there are exactly four combinations of positions for those two colours, because there are two cards for each colour.
- 2. The colours are the same: In each game there is only one combination of positions for those colours.




Figure 5.2: Exemplary showing the number of different comparisons for case 1 and 2. The numbers on the cards represent the case and the edges indicate unique comparisons. All cards are turned face up. A game without any obstacles is used only for the purpose of clarifying the different comparisons. All screenshots used for the actual calculation are taken from glare effect games.

First it is calculated how many screenshots are necessary for the first case. For two different coloured arbitrary but fixed cards there can be

$$C_1 = 14 \cdot 13 = 182$$

combinations of positions. The reason that it is not not $14 \cdot 14$ is that 14 comparisons of cards with themselves would be included. We formulate the condition that enough

screenshots need to be taken so that an arbitrary but fixed combination of positions for two different colours is included with a probability of 95%.

As there are four such combinations of positions in each game, the probability $P(A)$ to get a specific combination in a game is

$$P(A) = \frac{4}{182}$$

The counter probability $P(\neg A)$ is

$$P(\neg A) = 1 - \frac{4}{182} = \frac{178}{182}$$

The number of necessary screenshots to include a specific combination with 95% probability is calculated by solving

$$\begin{aligned} 1 - P(\neg A)^n &\geq 0.95 \\ 1 - \left(\frac{178}{182}\right)^n &\geq 0.95 \\ 1 &\geq 0.95 + \left(\frac{178}{182}\right)^n \\ \left(\frac{178}{182}\right)^n &\leq 0.05 \\ n &\geq \log_{\left(\frac{178}{182}\right)}(0.05) \\ n &\geq 134.8024 \end{aligned}$$

This means that about 135 Screenshots of games are needed in the first case. Now we perform the same calculation for the second case where there is only one combination of positions for the colours in each game. However, there are less combinations of positions that are relevant than in the first case. For example, for two green cards at index 0 and 1 the 182 comparisons would include a comparison of the first green card and the second green card as well as a comparison of the second green card and the first green card. This goes for every two cards with the same colour, meaning there are only half as many different comparisons in the second case than in the first case. As a result, the number of possible combinations in the second case is

$$C_2 = \frac{C_1}{2} = \frac{14 \cdot 13}{2} = 91.$$

Now we calculate how many screenshots must be taken to include a specific combination of positions for two equal colours with a probability of 95%.

As there is one such combinations of positions in each game, the probability $P(A)$ to get a specific combination in a game is

$$P(A) = \frac{1}{91}$$

The counter probability $P(\neg A)$ is

$$P(\neg A) = 1 - \frac{1}{91} = \frac{90}{91}.$$

The number of necessary screenshots to include a specific combination with 95% probability is calculated by solving

$$\begin{aligned}
 1 - P(\neg A)^n &\geq 0.95 \\
 1 - \left(\frac{90}{91}\right)^n &\geq 0.95 \\
 1 &\geq 0.95 + \left(\frac{90}{91}\right)^n \\
 \left(\frac{90}{91}\right)^n &\leq 0.05 \\
 n &\geq \log_{\left(\frac{90}{91}\right)}(0.05) \\
 n &\geq 271.111
 \end{aligned}$$

This means that about 272 screenshots of games are needed in the second case. As there are more screenshots needed for the second case, taking that many screenshots is also sufficient for the first case. This is caused by the fact that each screenshot includes comparisons for the first as well as the second case. To have a buffer the decision was made to take screenshots of 300 games. In theory these screenshots include any arbitrary but fixed combinations of positions and colours with a probability of over 95%, which should eliminate the problem of the positional influence.

Nonetheless, this does not solve the second problem of extracting single pixels that might not be representative of the overall perceived colours. This can be fixed by extracting all pixels in a certain area of the card and averaging their RGB values.

Last but not least it needs to be clarified how the colour differences are calculated. To capture how colour differences are perceived by humans, the RGB colour space is not suitable, as it is perceptually non uniform. A colour space is perceptually uniform if the measured colour differences are proportional to the human perception of such differences. Contrary to the RGB colour space, the L*a*b* colour space is perceptually uniform [Pas01, Chapter 1]. Therefore, the colour differences described in the similarity matrix are calculated after converting the colours into the L*a*b* colour space. In related work, when creating the similarity matrix for the effect of colour blindness, the CIE1976 colour model was used for the calculations [SPIS19, p. 5] of the colour differences. However, the formula has since then been improved multiple times, resulting in the CIE2000 colour model. Through multiple modifications of the formula the calculated colour differences became closer to how they are perceived by humans [Fis02, p. 12]. Therefore, when creating the similarity matrix for the glare effect the CIE2000 colour model is used instead of the CIE1976 colour model. The resulting colour distances are referred to as Delta E values [Fis02, p. 15]. The lower the Delta E value, the more similarly the colours are perceived by humans.

In summary, the chosen approach is to take 300 screenshots of glare effect games with all cards turned face up, extracting areas of RGB values for each card, converting the average RGB values of areas into L*a*b* colours, calculating the Delta E values for each combination of colour to create a similarity matrix for each screenshot and finally averaging the values in the 300 similarity matrices. As there are cards with seven different colours in the game, the similarity matrix has 28 entries (21 comparing non-matching and 7 comparing matching pairs). To achieve values between 0 and 1, the colour differences are additionally scaled down in the end.

5.1.2 Screenshot extraction

To play the memory game and take screenshots an emulator for the Pixel 2 in Android Studio was used. In order to take the screenshots in the needed way and collect additionally necessary information some changes to the memory game were made. In the way the game is intended there are always only at maximum two cards turned face up. However, for the screenshots all cards need to be turned face up so that the colour differences can be calculated. Therefore, the memory game was adjusted so that all cards are turned face up once the player turns a card. Additionally, it was changed so that cards stay face up once they get turned around for the first time. This makes it possible to take screenshots with the colours of all cards visible. Furthermore, for the creation of the similarity matrix it needs to be known what the original colours without the influence of the simulated sun are. To collect the screenshots and the according colours of the cards in each game a semi-automatic approach was chosen. Once a game is started and a card is flipped, resulting in all cards being turned face up, the names of the original colours of all cards are saved in a list. Then a screenshot is manually taken and afterwards the game is exited to enter the main menu. From there on the process is repeated 300 times. As a result, there are 300 screenshots of games and a two dimensional list that includes the colours of the cards in the 300 games. The first dimension specifies the game and the second dimension the index of the card. The values of this list are stored in a text file before the game is completely closed. To assure that no mistakes were made when taking the screenshots, the number of collected screenshots is observed during the whole process and afterwards all screenshots are manually checked to picture all cards turned face up.

5.1.3 Implementation

This section will show and explain the implementation of the similarity matrix generation. Code is only included if it helps to further understand what is being explained. Firstly, the screenshots and the information about the original colours of the cards are loaded. Before the actual calculation of a similarity matrix for each image, the average RGB values of the cards on the field need to be determined.

```

1  def determine_glare_rgb_values(image):
2      '''
3      Calculates the average rgb values for each card in an image.
4      :param image: The screenshot of the memory game with all cards
5      turned face up.
6      :return: The average rgb values in the specified 150 x 150 pixel
7      areas for each card.
8      '''
9      glare_rgb_values = []
10     for corner in card_corners:
11         x_border = corner[0] + 150
12         y_border = corner[1] + 150
13         card_values = []
14         for x in range(corner[0], x_border, 1):
15             for y in range(corner[1], y_border, 1):
16                 coordinates = x, y
17                 pixel_values = image.getpixel(coordinates)
18                 card_values.append(pixel_values[: -1])
19         card_r = int(round(np.mean([color[0] for
20             color in card_values])))
21         card_g = int(round(np.mean([color[1] for
22             color in card_values])))
23         card_b = int(round(np.mean([color[2] for
24             color in card_values])))
25         glare_rgb_values.append((card_r, card_g, card_b))
26     return glare_rgb_values

```

Listing 5.1: Code for calculating the average rgb values of each card in an image.

The cards are each around 250 x 250 pixels big and the RGB values are extracted from squared 150 x 150 pixel areas (given a resolution of 1920 x 1080). The coordinates the pixels are extracted from are manually selected in gimp. As a result, the areas are only correct for the used resolution of 1920 x 1080. Once the mean RGB values in those areas are calculated and converted into L*a*b* colours the similarity matrix can be created. L*a*b* colours are necessary to calculate the colour difference using the CIE2000 colour model. Each of the 28 cells in the similarity matrix falls into one of the two cases explained in section 5.1.1 [Conceptualization](#), meaning there are either 4 or only 1 unique combination for the calculation of each cell value in each screenshot. For every combination the Delta E value is determined and the values for each cell are averaged. This completes the steps for a single image, resulting in a similarity matrix with unscaled values. Once a similarity matrix with unscaled values for every image is created, the average of all matrices is calculated. During the whole calculation of the single matrices it is being kept track of the highest occurring colour difference. The last step necessary to complete the final similarity matrix is to divide all values in the matrix by the highest occurred colour difference in order to scale them down to values between 0 and 1.

5.1.4 Testing

To verify the correctness of the colour extraction and the calculation of the similarity matrix, the main functionalities are manually tested. The colour extraction was tested by extracting colours of a screenshot with the script and comparing the RGB values to the actual values in the images using gimp. Furthermore, the calculation of the Delta E score was tested, by comparing results of the script with those of an online calculator [Cie].

To assure that the 300 games actually include most of the combinations and therefore eliminate the positional influence of cards, the coverage of combinations can be additionally calculated. Therefore, the number of unique combinations included in the 300 screenshots is divided by the overall number of possible combinations. The number of unique combinations can be counted during the process of creating the similarity matrix, but the number of possible combinations must separately be determined. The similarity matrix for glare effect has 28 entries, from which 7 are comparisons between same and 21 between different colours. With C_1 and C_2 being the possible combinations for two arbitrary but fixed coloured cards in the two cases mentioned above, the number of overall possible combinations C is

$$\begin{aligned} C &= 21 \cdot C_1 + 7 \cdot C_2 \\ &= 21 \cdot 182 + 7 \cdot 91 \\ &= 4459. \end{aligned}$$

The 300 games used for creating this matrix include 99.57% of all possible combinations. The final test of the created similarity matrix, however, will be to actually use it to simulate glare effect games and see whether they are simulated accurately. This is done in section [5.6 Evaluation of simulation](#).

5.1.5 Final Matrix

In figure 5.3 the final similarity matrix can be seen, displaying how colour differences are observed under the influence of the simulated sunlight, averaged over 300 games. The maximum Delta E value that occurred during the calculation was 8.861. All values are divided by that number to achieve scores between 0 and 1. The lower the value the less different the cards are perceived by humans.




Figure 5.3: The final similarity matrix containing similarity values for each combination of coloured card on the field. The lighter the colour used for the representation the less of a difference is noticeable.

Prior to actually using the constructed similarity matrix in the simulator it is unknown if the chosen approach results in simulations of high quality. If this is not the case it is possible to change the concept or try other approaches.

It should be also noted that the similarity matrix has 28 entries. The one that was used for describing the confusion that is caused by a red-green colour vision deficiency only has 21 entries, because all 7 values describing cards of the same colour can be left out as the colours are equal. Because in glare effect games, originally equally coloured cards can look different due to the influence of the simulated sunlight, these 7 values must be included in the glare effect similarity matrix.

5.2 Incorporation of the new similarity matrix

The newly created similarity matrix can now be added to the glare effect logs. Therefore, each of the 28 similarity values in the matrix is assigned to a dynamically mapped card code. To do so there was already a script available that was used after some small adjustments. The main changes were to replace the deprecated functions from libraries with supported ones and incorporate the new similarity matrix. As the code was initially written for similarity matrices with 21 entries and the one for glare effect has 28 entries, the code needed further adjustments. After all changes were completed, the new glare effect logs with remapped logs were saved in new files. To check whether the similarity assignments are still done correctly, the resulting logs were compared to logs that were created by the original script. Besides different similarity values due to different matrices as well as the fact that the logs now also contained entries for same coloured cards, the mapping was identical.

5.3 Removal of invalid logs and balancing of the data set

For the simulation to work, the game log that is used for the simulation must have had all cards turned around at least once. If this is not the case the log is classified as invalid. Additionally to removing invalid logs, the aim is to balance the data set for training. Data imbalance can force the classification algorithm to be biased towards the majority class. This can lead to concepts of the minority class not being learned adequately [WLW⁺16, Chapter 1]. Therefore, equal numbers of games from each of the two game modes are included. Furthermore, the decision was made that the number of games from each participant in each game mode should be equal, too. This means that the data should not contain more no obstacle games from a participant than glare effect games, and vice versa.

As stated in chapter 4 [Collected Data](#), from the 22 participants there are 44 no obstacle and 21 glare effect game logs. To collect the data used for the simulation, a script was written that collects one no obstacle and one glare effect log from each participant. Half of the available no obstacle logs are not collected, because there are not as many glare effect logs. Once the logs are collected, they are validated and the valid ones are saved in new files. If at least one the two logs from a participant from different game modes is invalid, both are removed. Otherwise the training data would be inconsistent in that sense that a no obstacle log from a subject might be included but no glare effect log and vice versa. From the 22 no obstacle and 21 glare effect logs collected by the script, one glare effect log was invalid. This resulted in 20 no obstacle and 20 glare effect logs being used for simulation. These 40 real logs combined with those simulated form the data available for training.

5.4 Simulation of user behaviour

The whole process of simulating user behaviour is divided into two parts: First, configuration files are created using a generic optimization algorithm that contain the optimized parameter values and secondly, these configuration files are used to simulate game logs. How the simulator functions is explained in chapter 3 [CMM-based Cognitive User Simulation](#).

Multiple additions were made to the simulator. In total four classes were added. Two are for generating configuration files and the other two for using those files to simulate new game logs with and without the glare effect obstacle based on the user behaviour in the original games. These classes utilize the functionalities already implemented in the simulator. However, as the simulator only worked with similarity matrices that have 21 entries, instead of the 28 of the matrix for glare effect games, the simulator was adjusted so that it can also handle matrices with 28 entries. After all changes and additions were completed, each log was used to simulate 1000 logs, resulting in 40000 logs. From these logs 20000 are no obstacle logs and the other 20000 are glare effect logs. The 40000 logs contain 1000 no obstacle and 1000 glare effect logs for each of the 20 subjects.

5.5 Sorting logs by quality

The simulated logs were sorted by how close their performance is to that in the real game used for their simulation. The value by which is sorted is the average of two others. The first one is the root mean squared error (RMSE) between the matching pairs in the real and the simulated game for each round. The second one is the RMSE between the penalties in the real and the simulated game for each round. These two performance measurements are explained in section 5.6 [Evaluation of simulation](#). The formula for calculating the RMSE is shown below.

$$RMSE = \sqrt{\left(\frac{1}{n}\right) \sum_{i=1}^n (y_i - x_i)^2}$$

n : number of rounds
 y_i : performance value in real game
 x_i : performance value in simulated game

Sorting the logs according to their quality makes it possible to only use a certain number of the best simulated logs during the training instead of all of them or a random subset.

5.6 Evaluation of simulation

To evaluate whether the performance in the simulated logs is similar to the one in the original logs two performance measurements are used: The matching pairs in each round and the penalties in each round. Penalties are given if a card is revealed whose partner has already been seen before but the pair was not picked up. Some initial observations inspired changes to the simulator. The quality of the simulations regarding the performance measurements is analysed before and after the changes to the simulator by comparing the performance between simulations and real games. Furthermore, a subset of the simulated data was discovered that mostly shows no significant difference between the performance of the simulated and the real games.

5.6.1 Results before and after changing the simulator

The Initial simulation results can be seen in figure 5.4 and 5.5. The horizontal lines extending from the curves, also called whiskers, describe the standard deviation of the corresponding value.




Figure 5.4: Comparison of matching pairs per round of simulated and real glare effect and no obstacle games. Before the changes to the simulator.




Figure 5.5: Comparison of penalties per round of simulated and real glare effect and no obstacle games. Before the changes to the simulator.

At first glance it can be seen that the simulations of glare effect games are very similar to the real games, meaning that the constructed similarity matrix creates good simulation results. As a result, there is no need to find a different approach for creating the similarity matrix, and no changes to the simulator regarding the simulations of glare effect games are made. However, the simulation of no obstacle games is not quite as accurate. Especially the number of matching pairs after the tenth round in the simulated games is lower than in the real games. Additionally, the standard deviation indicated by the whiskers regarding the no obstacle games is noticeably higher in simulated than in the real games. This is not very strong in the first turns but gets worse in later turns. Furthermore, however less noticeable, are the penalties in the simulated games between round 8 and 18 lower than in the real games.

The fact that matching pairs per round for the no obstacle simulations are noticeably lower than in the real games inspired some changes to the simulator. As there are less cards on the field in later rounds and the experience of the user may have improved over the course of the game, in reality the degree of randomness is likely decreased in later turns. To simulate this a new parameter was introduced that is optimized by the generic optimization algorithm. It is called *randomizing decay* and is a value between 0 and 1. After the 10th round this value is used to reduce the amount of randomness in the choice of the next card and by that increases the performance in those turns. Another value that also influences the random behaviour in the simulations was optimized so that there is overall less randomness. The results of these changes can be seen in figure 5.6 and 5.7.




Figure 5.6: Comparison of matching pairs per round of simulated and real glare effect and no obstacle games. After the changes to the simulator.




Figure 5.7: Comparison of penalties per round of simulated and real glare effect and no obstacle games. After the changes to the simulator.

As the changes only affected the simulation of no obstacle logs, the following statements only refer to those. It is noticeable that the number of matching pairs per round in the simulated games is closer to the real games. Additionally, the standard deviation of the simulated games regarding the matching pairs per round gets noticeably smaller after the 10th round and by that is closer to that in the real games. However, the reduction of randomness also resulted in slightly less penalties after round 10 and by that the difference in penalties between the simulations and the real games becomes bigger. On the contrary, the standard deviation of the penalties in the last six rounds of the simulated no obstacle games got slightly closer to that in the real games. Reasons for the negative impact on the penalties are unknown. It seems that the improvements regarding the matching pairs overweight the deteriorations of the penalties, but if these changes actually improve the overall quality of the simulations can not be said without further analysis.

However, another interesting observation can be made. The standard deviations of the two performance measurements in the last rounds of the simulated no obstacle games is notably higher than in the real games. This might be caused by the fact that there are in general less or even no cards on the field if all pairs were discovered. This theory also fits to the fact that this is not an issue in the glare effect games, as there are more cards on the field in later turns. This concludes that there is a

possibility that the simulator struggles to accurately simulate rounds once there are only a few or no cards remaining on the field.

Although major tendencies of similarities and differences are visually noticeable in figure 5.4, 5.5, 5.6 and 5.7, a more accurate analysis is needed to make reliable statements. Therefore, multiple paired sample t-test are performed comparing real no obstacle, simulated no obstacle, real glare effect and simulated glare effect games with each other. For each combination two tests are performed. One comparing the mean penalties per round and the other one comparing the mean numbers of matching pairs per round. The p value expresses whether the compared lists of mean values are significantly different or not. Values higher than 0.05 indicate no significant difference, while values below 0.05 indicate a significant difference. It would be optimal if all comparisons between different game modes show significant differences, while comparisons between real and simulated games of the same mode show no significant differences. All paired sample t-tests were performed before and after the changes to the simulator. This was done to see whether the changes to the simulator improved or deteriorated the results of the test or if they had any impact on them at all. First the results of the paired sample t-tests, when including all 20 rounds and all simulated games, are analysed. Tables 5.1 and 5.2 show the results before changing the simulator, while tables 5.3 and 5.4 show the results afterwards.

The following abbreviations are used: real glare effect - r_g, real no obstacle - r_n, simulated glare effect - s_g, simulated no obstacle - s_n.

	s_n		s_g		r_n	
	p	t	p	t	p	t
r_g	$1.7e-06$	-6.8179	0.0688	-1.9292	$9.3e-07$	7.1041
r_n	$5.7e-07$	7.3571	$2.7e-06$	6.5765		
s_g	$5.3e-06$	6.2480				

Table 5.1: p and t values in paired sample t-tests for different comparisons of matching pairs per round before the changes to the simulator. All 1000 simulated games per real game were used.

	s_n		s_g		r_n	
	p	t	p	t	p	t
r_g	$3.6e-06$	-6.4297	$1.1e-06$	-7.0423	$3.2e-06$	6.4897
r_n	0.2011	-1.3242	$5.6e-06$	6.2238		
s_g	$5.4e-06$	6.2394				

Table 5.2: p and t values in paired sample t-tests for different comparisons of penalties per round before the changes to the simulator. All 1000 simulated games per real game were used.

	s_n		s_g		r_n	
	p	t	p	t	p	t
r_g	$5.7e-06$	-6.2181	0.0688	-1.9292	$9.3e-07$	7.1041
r_n	$1.8e-10$	12.2469	$2.7e-06$	6.5765		
s_g	$1.7e-05$	5.7107				

Table 5.3: p and t values in paired sample t-tests for different comparisons of matching pairs per round after the changes to the simulator. All 1000 simulated games per real game were used.

	s_n		s_g		r_n	
	p	t	p	t	p	t
r_g	$4.8e-06$	-6.2967	$1.1e-06$	-7.0423	$3.2e-06$	6.4897
r_n	0.0195	-2.5519	$5.6e-06$	6.2238		
s_g	$7.2e-06$	6.1045				

Table 5.4: p and t values in paired sample t-tests for different comparisons of penalties per round after the changes to the simulator. All 1000 simulated games per real game were used.

It can be seen that all tests comparing different game modes show significant difference in penalties per round as well as matching pairs per round, which is desired. However, significant difference is not desired in the four tests comparing real and simulated games of the same mode, which are highlighted by colours. Before the changes to the simulator, the matching pairs per round in simulated and real no obstacle games show significant differences. The same goes for the penalties per round in simulated and real glare effect games. As the changes to the simulator only affected the simulation of no obstacle games the significant difference between the matching pairs in simulated and real glare effect games prevails. Despite the fact that the matching pairs per round in simulated no obstacle games seem closer to the real games in figure 5.6, the statistical tests still show a significant difference between them. Additionally, the small deterioration of the penalties that were caused by the changes to the simulator and are observable in figure 5.7 led to a significant difference in the penalties per round between simulated and real no obstacle games. Before the changes to the simulator this difference was not significant.

In total, before the changes to the simulator 2 out of 4 tests comparing games of the same mode result in undesired significant differences. After the changes to the simulator this increased to 3 out of 4 tests. However, this does not lead to the conclusion that the changes to the simulator decreased the quality of the simulations, as the matching pairs per round in simulated no obstacle game shown in figure 5.6 got noticeably more accurate.

5.6.2 Using a subset of the data

In the statistical tests in section 5.6.1 *Results before and after changing the simulator*, all 1000 simulations per participant are used. As some simulations are better than others, using a subset only containing a number of best simulations could lead to further findings. Since the simulated games are sorted by their quality in section 5.5 *Sorting logs by quality*, it is possible to only use a certain number of best simulations for the comparisons, and see whether this improves the undesired properties mentioned above. Furthermore, the number of rounds in the comparison can be changed, so that not all 20 rounds are included. The consideration of including less rounds is of importance because it can reveal which parts of the game are simulated accurately and which are not. The search for an configuration of those two settings was manually done by varying the number of game rounds included and the ratio between simulated and real data. A configuration of those two settings that produces interesting results consisted of using only the best 15 simulated games per game mode from each participant and only including the first 10 rounds of each game. The results of the paired sample t-tests before the changes to the simulator can be seen in table 5.5 and 5.6, while the results afterwards are shown in table 5.7 and 5.8. The two performance measurements in the first 10 rounds after the changes to the simulator when only using the best 15 simulated games per game mode from each participant are additionally portrayed in figure 5.8 and 5.9

The following abbreviations are used: real glare effect - r_g, real no obstacle - r_n, simulated glare effect - s_g, simulated no obstacle - s_n

	s_n		s_g		r_n	
	p	t	p	t	p	t
r_g	0.0087	-3.3342	$6.4e-05$	6.9956	0.0058	3.5985
r_n	0.0032	3.9837	0.0037	3.8893		
s_g	0.0056	3.6218				

Table 5.5: p and t values in paired sample t-tests for different comparisons of matching pairs per round before the changes to the simulator. Only the best 15 simulated logs from each real log and the first 10 rounds are used.

	s_n		s_g		r_n	
	p	t	p	t	p	t
r_g	0.0025	-4.1496	0.1741	1.4756	0.0014	4.5577
r_n	0.6650	0.4476	0.0012	4.6627		
s_g	0.0021	4.2596				

Table 5.6: p and t values in paired sample t-tests for different comparisons of penalties per round before the changes to the simulator. Only the best 15 simulated logs from each real log and the first 10 rounds are used.

	s_n		s_g		r_n	
	p	t	p	t	p	t
r_g	0.0064	-3.5291	$6.4e-05$	6.9956	0.0058	3.5985
r_n	0.0763	3.8893	0.0037	3.8893		
s_g	0.0041	3.8201				

Table 5.7: p and t values in paired sample t-test for different comparisons of matching pairs per round after the changes to the simulator. Only the best 15 simulated logs from each real log and the first 10 rounds are used.

	s_n		s_g		r_n	
	p	t	p	t	p	t
r_g	0.0019	-4.3207	0.1741	1.4756	0.0014	4.5577
r_n	0.8388	-0.2095	0.0012	4.6627		
s_g	0.0016	4.4302				

Table 5.8: p and t values in paired sample t-test for different comparisons of penalties per round after the changes to the simulator. Only the best 15 simulated logs from each real log and the first 10 rounds are used.

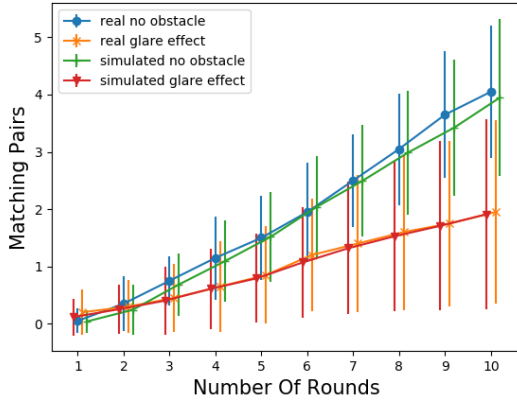


Figure 5.8: Comparison of matching pairs per round of simulated and real glare effect and no obstacle games in the first 10 rounds. Only the best 15 simulations per real game were used. Furthermore, the data used is from after the changes to the simulator.




Figure 5.9: Comparison of penalties per round of simulated and real glare effect and no obstacle games in the first 10 rounds. Only the best 15 simulations per real game were used. Furthermore, the data used is from after the changes to the simulator.

Interestingly, including all the simulated data in the tests, results in more undesired significant differences before the changes to the simulator compared to afterwards, but using only a small subset of the data results in less. Before the changes to the simulator, when including only the best 15 games and the first 10 rounds in the tests, only the two 2 tests in table 5.6 for penalties comparing real and simulated games of the same mode show no significant difference. The two tests in table 5.5 comparing matching pairs of real and simulated games of the same mode show a significant difference. Contrary to that, after the changes to the simulator, the test in table 5.7 for matching pairs in real and simulated no obstacle games shows no significant difference, meaning more tests have the desired outcome.

As all but one test produce desired results after the changes to the simulator, it was additionally searched for a configuration that only result in desired outcomes of the paired sample t-tests. The only undesired result after the changes and using the specified subset of the data is highlighted in red in table 5.7. It expresses a significant difference in matching pairs between real and simulated glare effect games. However, a configuration that showed no significant difference in all tests were it is desired was not found. To achieve non-significant difference between matching pairs in real and simulated glare effect games, a noticeably larger subset of the data had to be included, which resulted in a significant difference in matching pairs between real and simulated no obstacle games. Nonetheless it should be pointed out that this does not mean that the simulation of the glare effect games is bad regarding the number of matching pairs per round. If the two lists of mean values for matching that are significantly different according to the paired sample t-test are manually compared, it can be seen that the difference is very small, with the highest difference being 0.093 in round 6. These values can be seen in table 5.9.

	r. 1	r. 2	r. 3	r. 4	r. 5	r. 6	r. 7	r. 8	r. 9	r. 10
real	0.2	0.3	0.45	0.65	0.85	1.2	1.4	1.6	1.75	1.95
simulated	0.117	0.253	0.403	0.61	0.793	1.107	1.32	1.527	1.71	1.91

Table 5.9: Mean values of matching pairs for real and simulated glare effect games for the first 10 rounds. Only the 15 best simulations per real game are used.

5.6.3 Conclusion of the evaluation

Firstly, it can be said that using the created similarity matrix for the glare effect simulations produces games with performance close to the real games. Therefore, the aim of creating a similarity matrix that describes the colour differences of cards on the field under the influence of the glare effect was achieved.

Regarding whether the attempt to improve the simulation of no obstacle games by doing changes to the simulator was successful, no clear statement can be made. When all simulated data is used, the paired sample t-tests result in more undesired outcomes after the changes than before, while using only the best 15 simulations and only the first 10 rounds results in less undesired outcomes after the changes compared to before.

Although no subset of the data was found in which all simulations have no significant difference to the according real games regarding the paired sample t-test, it can be said that the simulator is capable to simulate the first 10 rounds very accurately, given only the 15 best simulations are used. The statistical tests and the previous considerations regarding the too large standard deviation in later turns indicate, that the simulator performs less good in later turns. Having said that, these are only indications that would need further research to be confirmed. Doing so might lead to improving the simulator.

This analysis also indicates that it might not ideal to use all simulated data for training as that results in a significant difference between the simulated and the real games according to the paired sample t-test. Using too many simulations compared to real games could decrease the performance of the classifiers on real data if they adapt too much to simulated games. Furthermore, it should be analysed how using different amounts of turns in the classification impacts the classification results. In section 8.3 Training results and evaluation different configurations regarding ratios of simulated to real games and the number of rounds included, are tested.

Nonetheless it should be emphasized that the aim of no significant difference between real and simulated games is set very high. As shown in table 5.9, the difference can be very little and the paired sample t-test still concludes a significant difference. The fact that some tests show a significant difference might not be related to bad simulations but to the limited number of pairs used (7 pairs, 14 cards) [SPIS19, p. 6]. Overall the simulations and real games are very similar when comparing them with the two performance measurements. If only the best simulations are used the performance becomes even more similar.

5.7 Feature generation

5.7.1 1D convnet features

For the 1D convnet five statistical features for each step (flipping of a card) are used: The card codes, the number of cards left, the number of never revealed cards, the highest number of times the same card was revealed and the number of steps since all pairs were found (one round consist of two steps). They are calculated using the original game logs as well as the ones simulated in section [5.4 Simulation of user behaviour](#). Their structure is explained in chapter [4 Collected Data](#). The calculated features can be directly fed into the 1D convnet constructed in section [8.1.1 1D convnet](#).

The code for calculating the statistical features out of game logs was already given. A script was written that incorporates this functionality in order to calculate the features for all logs. As mentioned in section [8.1 Models](#), leave-one-out cross validation is used, where each model is trained with data from $n-1$ subjects and thus tested on the data from the remaining subject [[Won15](#), Chapter 1]. With 20 games from each of the to relevant game modes this results in 20 splits. After creating the necessary directory structure the script saves the files for the splits of the raw data and the features. The reason for saving the features is that they do not have to be calculated before every training and instead can be loaded from files. The reason for also saving the raw data in splits even though this work does not need them is, that the working group that was collaborated with also trained other models and does not directly load the features but instead calculates them before each training.

5.7.2 2D convnet features

For the second approach of using a 2D convnet further steps are taken. Synthetic images in grey scale colours are created using the features mentioned above. As the images are in grey scale colours, only one colour channel is needed. One can visually think of this approach as creating graphs for each feature that display their values in each step, taking a photograph of each graph and stacking them on top of each other. This can be clarified by looking at the code that produces the synthetic images.

```

1  def create_image(game, components=[True, True, True, True, True]):
2      '''
3      Creates synthetic images out of the card codes and the staticial
4          features.
5      :param game: The statical features in each step.
6      :param components: Five values describing which of the features
7          should be used to create the image.
8      :return: The synthetic image.
9      '''
10     card_codes = np.zeros((7, steps))
11     cards_left = np.zeros((8, steps))
12     never_revealed_cards = np.zeros((14, steps))
13     max_same_card_reveals = np.zeros((20, steps))
14     rounds_since_done = np.zeros((27, steps))
15
16     x_position = 0
17     for step in game:
18         card_code = math.floor(step[0])
19         first_or_second = int(round((step[0] % 1) * 10))
20         if card_code != 0:
21             card_codes[card_code - 1][x_position] = first_or_second
22         pairs_left[int(step[1] / 2)][x_position] = 1
23         never_revealed_cards[int(step[2])][x_position] = 1
24         max_same_card_reveals[int(step[3])][x_position] = 1
25         rounds_since_done[int(step[4])][x_position] = 1
26         x_position += 1
27
28     image = np.zeros((0, steps))
29     if components[0]:
30         image = np.vstack((image, card_codes))
31     if components[1]:
32         image = np.vstack((image, max_same_card_reveals))
33     if components[2]:
34         image = np.vstack((image, rounds_since_done))
35     if components[3]:
36         image = np.vstack((image, cards_left))
37     if components[4]:
38         image = np.vstack((image, never_revealed_cards))
39     return image

```

Listing 5.2: Code for creating a synthetic image from the statistical features.

By using pseudo colors, the images can be displayed with colours, like in figure 5.10. These 76 x 40 x 1 (height x width x colour channels) images are directly fed to the 2D convnet constructed in section 8.1.2 2D convnet. By setting the origin of the image to the lower instead of the upper left corner a more natural looking image is created and can be seen in figure 5.11.




Figure 5.10: Synthetic image created using the statistical features. Origin of the image in the top left corner.

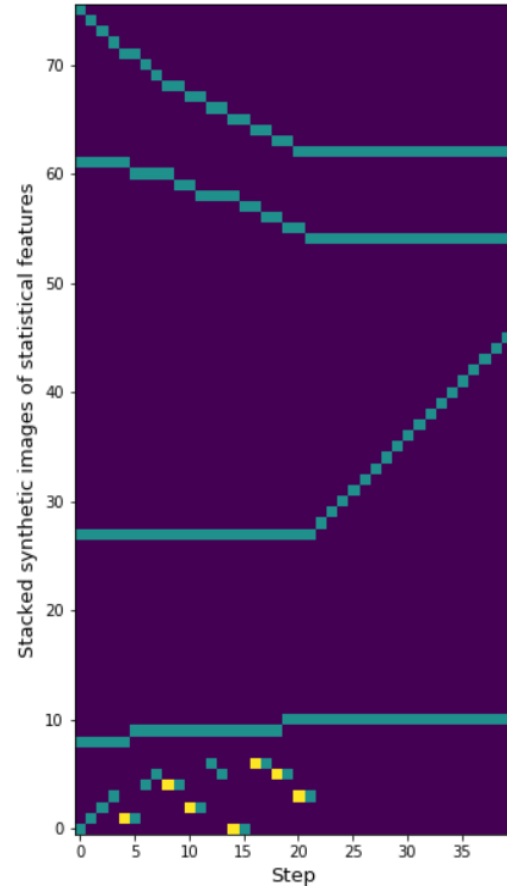


Figure 5.11: Synthetic image created using the statistical features. Origin of the image in the bottom left corner.

The width of the image is 40 because that the number of steps recorded. Table 5.10 shows which of the areas in the image describe which statical feature. The different vertical spaces given to statistical features were purposefully chosen. The aim was to use as much space as necessary but as little as possible. As each round consists of two cards being turned face up and because of the impossibility to choose the same card twice in one turn, a card can be at maximum revealed 20 times during 20 rounds. Considering that this value is at minimum 1 because it is not possible to flip no cards, the range of 20 values is sufficient for describing the maximum of same card reveals. Furthermore, 14 steps are at minimum needed to complete the game, since there are 14 cards. As a result, the steps since the game ended can range from 0 to 26, meaning a vertical space of 27 is needed. In order to save space, the statistical feature of the number of cards left was converted to the number of pairs left. By dividing by 2 the vertical space necessary to visualize this feature is halved, without information being lost. Instead of representing the numbers of cards left

on the field, the value now represents the number of pairs left. Furthermore, the number of never revealed cards can range from 0 to 13. The value 14 is not possible because two cards have to be chosen each turn, meaning that it is impossible to not reveal any card. The stacking order was chosen so that the coloured pixels of different statistical features rarely touch each other.

Statistical feature	Range	Vertical space in the image
Card codes	1-7	0-6
Maximum of same card reveals	1-20	7-26
Steps since game ended	0-26	27-53
Pairs left	0-7	54-61
Never revealed cards	0-13	62-75

Table 5.10: Range and vertical space of the different component in the synthetic image.

Regarding the visual representation of the card codes two decisions were made: Using different colours for the first and the second card of a colour and combining the information about the 14 cards in an vertical area of size 7. For each card code the vertical position of the representing pixel is decided by the first number and the colour is chosen according to the second number. As a result, each row contains the information about cards of a specific colour. This representation of the card codes is not chosen arbitrarily. On the contrary, it was inspired by reasons that need explanation.

Instead of combining the information about the 14 cards in an vertical area of size 7, it would have been also possible to use double the vertical space so that the first and the second card of a colour each have separate areas. The first cards and the second cards of a colour would each have individual vertical spaces with sizes of 7. However, it was decided against it. This decision was based on a consideration regarding the structure of the convolutional neural networks used in this work and the concept of the receptive field, explained in section 6.3 [Receptive field](#). How convnets function is explained in chapter 6 [Convolutional neural networks](#) and which models are used in this work is described in section 8.1 [Models](#). In short, the structure of the model influences the sizes of the receptive fields of the features in layers of the network. The receptive field of a feature describes how much of the original input image has influence of the feature. The size of the receptive field is generally bigger the more convolution and pooling layers the network consists of. This also means that if the structure of the convnet is chosen small, the receptive field size of individual features might not be large enough to be influenced by a local area big enough to include far distant pixels of the input image. In this work, the 2D convnet used is constructed with a small number of layers. Therefore, it is desirable that pixels in between an important dependency can be assumed are close to each other so that the small convnet is able to learn that dependency. Typically in classical image recognition it is only possible to choose the structure of the models depending on what is to be recognized in the image. However, since in this case the images used are created synthetically, it is also possible to create them while considering the structure of the model. Meaning the images can be created so that they are better suited for the models. The chosen representation of the card codes exploits exactly this freedom. As the dependencies between cards of the same colours can be estimated to be very

important for the classification, the representation is chosen so that if cards of the same colour are flipped in quick succession the pixels describing them are also locally close to each other in the synthetic image. The dependency between such close pixels does not require a large receptive field in order to be learned by individual features of the network. Therefore, this representation is better suited for the small structure of the convnet that is being utilized.

The main reason that motivated the use of different colours derived from the first decision to combine information about 14 cards in a vertical space of 7. If the same colour was used for all card codes, there would be no visual difference between flipping the same card in consecutive turns and flipping two different cards with the same colour. This of course would also mean that the model could not differentiate between those two cases, although they stem from completely different behaviour. Using different colours for the first and the second card of a colour fixes this issue.

These two decisions resulted in an interesting visual difference between no obstacle and glare effect games that could potentially be beneficial if learned by the model. By comparing the the card code sections of two images from which one is created using a glare effect and the other one with a no obstacle game, a noticeably visual difference can be observed.




Figure 5.12: Card code section of a synthetic image for a no obstacle game.




Figure 5.13: Card code section of a synthetic image for a glare effect game.

In 5.12, showing the card code visualisation for a no obstacle game, yellow pixels are directly left of green ones, meaning that the subject flipped a card, knew that he had already seen the matching card and directly flipped it. However, this happens less often in figure 5.13, showing card code visualization for a glare effect game. This is not the case in every game, but the theory is that separations of different coloured cells in the image are more likely in glare effect games than in no obstacle games. Therefore, it could be beneficial for the model to learn these characteristics. It should be noted that what the convnet exactly learns or what exact influence the chosen representation has can only be speculated. Although the way convnets function can be considered when deciding how to create the synthetic images, this does not guarantee good classification results.

6. Convolutional neural networks

This chapter introduces convolutional neural networks, also known as convnets. There exist different types that can be used for various tasks. While one dimensional convolutional neural networks (1D convnets) can be used for sequence classifications [Cho17, Chapter 6], two dimensional convolutional neural networks (2D convnets) are almost universally used in the area of computer vision [Cho17, Chapter 5]. There are other types of convnets, but this work utilizes only those two. Therefore, only they will be explained. Furthermore, only the concepts relevant for this work will be covered. To explain the concepts of convnets, 2D convnets will be used. Once understood, they can easily be applied to 1D convnets, as there is little difference between the two types of models. The differences that exist, will be explained.

Convnets consist of multiple components called layers [Cho17, Chapter 3]. Which combination of layers to use is highly dependent on the data and its complexity. Same goes for the choice of the so-called hyper parameters. These parameters play a major role in how the model is trained and how it will perform. Two important hyper parameters are the number of epochs and the learning rate [DHK13, Chapter 4]. While the number of epochs describes how often the whole training data should be passed to the network, the learning rate controls how much the network adapts to the training data in each epoch. How neural networks are trained is explained in chapter 7 [Training neural networks](#). More specific to convnets are the parameters for the size of the patches extracted from the inputs and the depth of the output feature map [Cho17, Chapter 5], described in section 6.1.1 [Convolution layers](#). Finding the best hyper parameters to optimize the network can be done in various ways, for instance using evolutionary algorithms [YRK⁺15]. However, this is not part of this work.

6.1 Layers

The convolutional neural networks used in this work consists of the following types of layers: convolution layers, dense layers, also known as densely connected layers, dropout layers, pooling layers and flatten layers. They will be explained without going too much into detail about the exact calculations performed.

6.1.1 Convolution layers

Convolutions operate over three dimensional tensors and are defined by two key parameters: The size of the windows that extract patches from the inputs and the depth of the output feature map. Typically the size of these windows is 3×3 or 5×5 . The depth of the output feature map can also be interpreted as the number of filters computed by the convolution. A convolution is performed by moving these two dimensional windows step by step across the three dimensional input feature map and extracting a three dimensional patch of surrounding features in each step. Each resulting patch has the height and width of the window and the depth of the input feature map [Cho17, Chapter 5]. Figure 6.1 exemplary shows the valid locations of 3×3 windows in a three dimensional input feature map with a height and width of 5 and an arbitrary depth.




Figure 6.1: Valid locations of 3×3 windows in a feature map with arbitrary depth and a height of 5×5 . Top view in the feature map. The depth is not shown. Source: [Cho17, Chapter 5].

Afterwards, each of these three dimensional patches is transformed into a one dimensional vector that has the specified depth of the output feature map. This transformation is done by calculating a tensor product with a weight matrix, called the convolutional kernel. The kernel has the same dimensions as the window that extracts the patches. The weights of the convolutional kernel are randomly initialized and learned during the training of the network. Once a vector for each patch is calculated they are combined into a three dimensional output feature map. This is done in a way so that every spatial location in the output feature map correspond to the same location in the input feature map [Cho17, Chapter 5]. The whole process of a convolution is shown in figure 6.2.




Figure 6.2: The whole process of a two dimensional convolution. Source: [Cho17, Chapter 5].

The way the convolution is performed is also what makes convolution layers fundamentally different from densely connected layers that are explained in section 6.1.5 [Dense Layers](#). Dense layers learn global pattern in their input feature space, while convolution layers learn local patterns. These local patterns are found in the patches of specified size extracted from the input. This characteristic gives convolutional neural networks two important properties. First of all the patterns they learn are translation invariant, meaning that for instance after learning a pattern in the lower left corner of an image it can be recognized anywhere in the image. Contrary to that densely connected networks would have to learn the pattern new if it appeared in a different location on the image. The second characteristic is that convolutional neural networks can learn spatial hierarchies of patterns. The first convolution layer can learn small local patterns like edges and the next convolution layer will receive the features of the first layer and learn larger patterns, and so on, resulting in the ability to learn increasingly complex patterns [Cho17, Chapter 5].

Last but not least it should be mentioned that the 3D input tensor of such a convolution layer can be created from an image, where the height and the width of the image are that of the tensor and the different colour channels of the image are represented by the depth of the tensor. However, after an image is passed through a convolution layer a new output feature map is produced and its depth is specified by one of the parameters mentioned above. As a result the depth channels no longer represent the different colour channels of the input image. Instead they are referred to as filters, that encode specific aspects of the input data [Cho17, Chapter 5].

6.1.2 Max-Pooling layers

There are multiple types of pooling layers, but max-pooling is the only pooling operation used in this work as it tends to work better than the alternative solutions [Cho17, Chapter 5]. Pooling operations are generally used to downsample the feature maps. Max pooling is performed by extracting windows, usually the size of 2×2 , from the input feature map and outputting the highest value of each channel. Such a window size halves the height and width of the feature map. This aggressive downsampling has two effects. Firstly, the computational complexity is reduced massively as there are less values to process and secondly, it has as a result that the data in the patches extracted from the inputs in successive convolution layers contain information about increasingly larger areas of the original input image. Therefore, the size of the region in the input data that has influence on the individual features created by the convolution is increased [Cho17, Chapter 5]. This increase is a fundamental concept of convolutional neural networks and is further explained in section 6.3 [Receptive field](#).

6.1.3 Dropout layers

The overall aim is to create models that generalize, meaning they perform well on data that they have never seen before. The central obstacle of generalization is overfitting. It describes the scenario in which the model adapts too much to the training data which reduced its ability to generalize. An effective way to reduce overfitting is using a dropout layer. Dropout is a technique originally developed by Geoff Hinton and his students at the University of Toronto [SHK⁺14]. It is applied to a layer by randomly setting a number of output features of the layer during training to 0. The rate in which this happens can be specified and is usually between 0.2 and 0.5. In short, the idea is, that randomly setting output to 0 produces noise that can break up random patterns which are not significant and otherwise would be memorized by the model if no noise was present [Cho17, Chapter 4].

6.1.4 Flatten layers

A flatten layer simply receives a multi dimensional input tensor, in case of 2D convnets this tensor is three dimensional, and flattens it into a 1D tensor that can be fed to a densely connected network [Cho17, Chapter 5].

6.1.5 Dense Layers

After flattening, the resulting 1D tensor is generally passed to a densely connected network [Cho17, Chapter 5]. This network consists of multiple dense layers that classify the input features. In dense layers each feature depends on the whole of the input data. Each layer applies tensor operations, that involve weights randomly initialized and learned during the training of the network. The final layer in a classification task often is a dense layer with softmax activation that returns an array in the size of the number of classes, specifying probabilities for each class [Cho17, Chapter 2]. Activation functions are briefly explained in section 6.2 [Activation functions](#).

6.2 Activation functions

In this work, activation functions are used in combination with convolution and dense layers. One activation function, the softmax function, has been already mentioned in section 6.1.5 [Dense Layers](#). It produces an array that implies the prediction of the model. In short, activation functions simply receive an input and calculate an output depending on that function. The other activation function that is utilized in this work is the rectified linear unit (relu). It enables the the layer to learn non-linear transformations of the input data. In comparison to otherwise being limited to linear transformations this gives the model access to a much richer hypothesis space. All the relu does is take the input and return 0 if the input is less than 0 and otherwise return the input value [Cho17, Chapter 3]. For instance, if combined with a convolution layer, each value in the output feature map is passed to the relu. This results in a new feature map that is passed to the next layer.

6.3 Receptive field

One of the fundamental concepts of convolutional neural networks is the receptive field, or field of view, of a feature in a certain layer of the network [LLUZ17, p. 1]. As mentioned before, in dense layers the features depend on the whole input, while a feature in a convolution only depends on a local region of the input. This region is the receptive field of that feature.

The receptive field size of features in the network can be increased in multiple ways. One way is by stacking more convolution layers, which increases the receptive field size linearly by theory. This effect is visualized in figure 6.3. Another, more effective way, is to use a pooling layer for subsampling, which increases the receptive field size multiplicatively [LLUZ17, p. 1]. As a result these methods can be used to increase the region size of the original input image that has influence on individual features in deeper layers.




Figure 6.3: Visualization of a linearly growing receptive field in stacked convolutions.

Source: [Ada20].

These considerations are important when building a model. As stated by Luo et al.: „Since anywhere in an input image outside the receptive field of a unit does not affect the value of that unit, it is necessary to carefully control the receptive field, to ensure that it covers the entire relevant image region“ [LLUZ17, p. 1]. The features are here referred to as units. This means that if there are regions of a certain size in the input image that should influence individual features as a whole, it is important to make sure when constructing the model that the receptive field in deeper layers is large enough for the specific task.

6.4 Differences of one dimensional convolutional neural networks

1D convnets can be used to classify sequence data. Instead of a 3D tensor, they receive 2D tensors as inputs. Such sequence data could for instance consist of muscle activity over time and the task could be to classify the type of movement. The convolutions explained in section 6.1.1 [Convolution layers](#) are two dimensional as the window that extracts patches from the image moves in two dimensions. The convolutions in 1D convnets are one dimensional, meaning the windows only move in one dimension and extract local patches from the sequence. Apart from this the convolution is performed similarly [Cho17, Chapter 6]. The process of a one dimensional convolution is shown in figure 6.4.




Figure 6.4: The whole process of a one dimensional convolution.
Source: [Cho17, Chapter 6].

The other difference is how the max-pooling operation is performed. Instead of extracting 2D patches and outputting the maximum value, this is done for 1D patches. The effects of this operation are the same as in 2D convnets [Cho17, Chapter 6].

The special characteristics of 2D convnets also apply to 1D convnets. They are translation invariant and capable to learn hierarchical patterns. Furthermore the dropout, flatten and dense layers as well as activations function in the same way. The behaviour explained in section 6.3 [Receptive field](#) can also be applied to 1D convnets: The patches extracted from the inputs in successive convolution layers contain information about increasingly larger areas of the original input sequence [Cho17, Chapter 6].

7. Training neural networks

Although the different components of convnets have been covered in chapter 6 [Convolutional neural networks](#), it has not been addressed yet, how these networks are trained to make accurate predictions. Therefore, the general concepts involved in training neural networks will be briefly explained.

The relationship between the core concepts revolving around training neural networks is shown in figure 7.1. The fundamental data structure in neural networks is the so called layer. There are many different types of layers that are chained together to form the network. In a classification task, what the network does is map the input data to a prediction. This is also called forward passing, as each layer performs operations on the received data and passes the results to the next layer. Once this chain reaches the output layer and by that outputs a prediction, the loss function takes that prediction as well as the actual target label and calculates a loss value which is a measure of how well the network's predictions match the expected outcome.

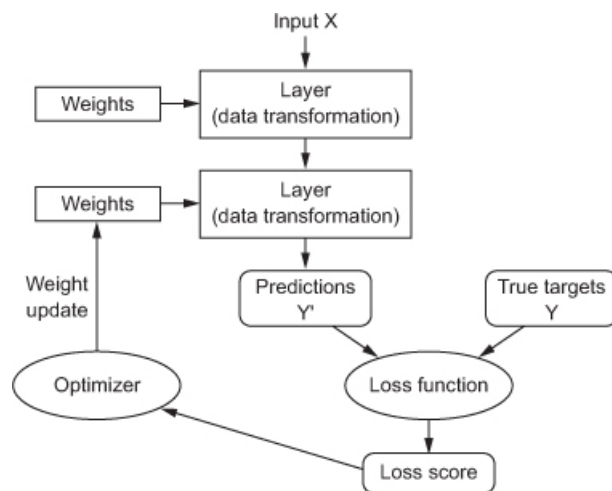


Figure 7.1: Core concepts in training neural networks. Source: [\[Cho17, Chapter 3\]](#).

This value is used by the optimizer to update the networks weights. This is done through an algorithm called backpropagation which works backward from the output layer to the input layer and computes the contribution that each weight had in the loss value. Using a method called stochastic gradient descent, the weights are adjusted in small steps to reduce the loss. The procedure of making predictions, calculating the loss and updating the weights is repeated many times, so that the loss is step by step reduced, resulting in more accurate predictions [\[Cho17, Chapter 3\]](#).

8. Classification

8.1 Models

For determining whether subjects played a game without any obstacles or a game effected by the glare effect, one dimensional and two dimensional convolutional neural networks are utilized. As the task at hand is naturally a sequence classification, using 1D convnets is the more conventional approach of the two. 2D convnets on the other hand are almost universally used in computer vision applications [Cho17, Chapter 5]. However, a possibility is to create synthetic images from the sequential data and use these images in a 2D convnet [Amu]. The synthetic images are created in section 5.7.2 2D convnet features. The dimensions of the data and the feature maps in figure 8.1 and 8.2 are calculated for the case that all 20 rounds are used. If a different number of rounds is used the values differ.

The models do not necessarily produce the best results possible since best practice 1D convnet [To] and 2D convnet [Amu] topologies are used. As the amount of collected real data with 20 recordings for each label is very limited, it was decided not to split the data into training, test and validation data, and instead only into training and test data using leave-one-out cross validation, where each model is trained with data from n-1 subjects and thus tested on the data from the remaining subject [Won15, Chapter 1]. As there is no separate validation set, an optimization would have to be done on the test data. If the parameters are optimized using the test set, the measurement of generalization would be flawed [Cho17, Chapter 4]. Therefore, it was avoided to optimize the hyper parameters individually for each model to ensure generalizing on new testing data. Instead, common best practice hyper parameters were optimized over all the models for better generalizing ability.

8.1.1 1D convnet

The structure of the 1D convnet used is shown in figure 8.1.




Figure 8.1: Visualization of the utilized 1D convnet topology. The dimensions represent those of the feature map in that layer. For example, the first convolution layer produces a feature map with the size 36 x 32 with 32 being the number of filters computed. The input is on the left and the output on the right.

The input data consist of the 5 statistical features, explained in section 5.7.1 1D convnet features, for each of the 40 steps (20 rounds), resulting in the input dimensions 40 x 5. Initially the data is passed to the first convolution layer of the network. The first layer computes 32 filters and the convolutional kernel has a size of 5, meaning that every step from the kernel includes all data from 5 steps in the game. The resulting feature map has a size of 36 x 32. This feature map is passed to the second convolution layer that computes the same number of filters. The kernel also has the same size like in the first layer. This results in a feature map with a size of 32 x 32. Afterwards, a dropout layer with a rate of 0.5 randomly sets input features to 0 to reduce overfitting. Then a one dimensional max-pooling layer with a pool size of 2 reduces the size of the feature map to 16 x 32. Afterwards, a flatten layer forms the features into one dimension. These features are finally passed to a densely connected network, consisting of two dense layers. Both convolution layers, and the first dense layer use a rectified linear unit as activation function. The last dense layer uses softmax as activation function to bring the data down to two values that specify the result of the classification. In total 57486 parameters are trained in this model.

For the optimization, the Adam algorithm is used that is already implemented in Keras. Adam optimization is a stochastic gradient descent method [Co18]. The learning rate is set to 0.0001 and not changed during the training. The batch size is set to 32 and the model is trained for 2500 epochs.

8.1.2 2D convnet

The structure of the 2D convnet used is shown in figure 8.2.




Figure 8.2: Visualization of the utilized 2D convnet topology. The dimensions represent those of the feature map in that layer. For example, the first convolution layer produces a feature map with the size $72 \times 36 \times 64$ with 64 describing the number of filters computed. The input is on the left and the output on the right.

The input data consist of the synthetic images generated in section 5.7.2 2D convnet features, using the 5 statistical features. This results in the input dimensions $76 \times 40 \times 1$. Initially the data is passed to the first convolution layer of the network. It computes 64 filters and the convolutional kernel has a size of 5×5 , meaning that each step from the kernel includes a 5×5 sized area of the synthetic image. This convolution results in a feature map with a size of $72 \times 36 \times 64$. This feature map is passed to the second convolution layer that computes 32 filters with a kernel of size of 3×3 . This results in a feature map with the size of $70 \times 34 \times 32$. Then a two dimensional max pooling layer with a pool size of 2×2 reduces the size of the feature map to $35 \times 17 \times 32$. Afterwards, a dropout layer with a rate of 0.2 randomly sets input units to 0 to reduce overfitting. Then a flatten layer forms the features into one dimension. These features are finally passed to a densely connected network, consisting of three dense layers. Both convolution layers, and the first two dense layers use a rectified linear unit as activation function. The last dense layer uses a softmax activation function to bring the data down to two values that specify the result of the classification. In total 2463928 parameters are trained in this model.

For the optimization, the Adam algorithm is used that is already implemented in Keras. Adam optimization is a stochastic gradient descent method [Co18]. The learning rate is set to 0.00001 and not changed during the training. The batch size is set to 32 and the model was initially trained for 500 epochs. Later, many 2D convnets had to be retrained due to complications and the epochs were reduced to 200, because 500 was unnecessarily high (accuracy continually decreased already after less than 500 epochs). This can be seen in figure 8.7.

8.2 Training setup

There is one script each for training the 1D convnet and the 2D convnet. The only differences are in the structure and the hyper parameters of the convnets and that before training the 2D convnet the loaded statistical features must first be used to create the synthetic images. Both scripts load the same statistical features, already divided into 20 splits that are used for leave-one-out cross validation. Each split contains a training and a test set. A training set consists of 19 real and 19000 simulated games for each of the two classes, while a test set contains one real game for each of the classes. The 2000 simulated games excluded from the training are not used for testing, because the results would not be representative of the real world performance of the models. Furthermore, only a certain number of the best simulations are included in the training. The ratio of n simulated games for each real games will be abbreviated with sdn x. For instance, $sd5$ x means that 5 simulated games for each real game were included in the training. Each training of a model is done for the following ratios: $sd0$ x, $sd1$ x, $sd2$ x, $sd3$ x, $sd4$ x, $sd5$ x, $sd6$ x, $sd7$ x, $sd8$ x, $sd9$ x, $sd10$ x and $sd20$ x. Moreover, the training is done on different numbers of rounds (each round consists of flipping two cards). In real world scenarios, the less turns waited the faster the decision is made and assistance can be applied. However, if waited longer the decision will likely be correct more often. To analyse this, the models are trained each for 5, 10, 15 and 20 rounds. More than 20 turns were not recorded. As convnets are stochastic machine learning algorithms, different runs of the same algorithm on the same data can result in different outcomes. Therefore, it is not sufficient to train each model only once as the result is not reliable. Instead, the training of a model is repeated multiple times and the results are averaged to summarize the performance of the model [Bro]. The decision was made to repeat the training of each model 20 times.

Once one of the scripts is executed, required directory structures are created and training processes for the different ratios of real to simulated games are performed successively. The training for one ratio consists of parallelized training of 20 splits using multiprocessing. As the training for each split is repeated 20 times, for each ratio 20 models are created for every split. This means that 400 models are trained for every ratio. As there are 12 different ratios trained, this results in 4800 models trained for each execution of one of the scripts. As mentioned above, the training is done for 4 different amounts of rounds for the 1D convnet and the 2D convnet. This adds up to 38400 models being trained. Furthermore, this is done twice: Once before the changes to the simulator and once afterwards, to see if the changes improved the performance of the models. All in all, this results in 76800 models. For each model the training history is saved in a file. Saving and loading the histories is necessary due to the parallel training of multiple models through multiprocessing. Once the training of the 4800 models of one script is completed, the training histories are loaded from the files. The results from the 400 training histories for each ratio of real to simulated games are averaged, plotted and the final figures are saved for the analysis. All models are trained on the CPU. The fact that 76800 models are trained led to the decision, not to save all models but instead first evaluate the results and then retrain and save the models that create the best results thus using less memory.

8.3 Training results and evaluation

The training results for the various configurations can be seen in table 8.1, 8.2, 8.3 and 8.4. The first two tables contain the results before the changes to the simulator and the last two the results afterwards. The analysis will only focus on the accuracy, but the loss is also listed for the sake of completeness. The best accuracies for each number of turns in each table are highlighted in green. If multiple ratios of real to simulated games for a certain number of turns produce the same accuracy only the one with the smallest amount of simulations is highlighted. The accuracies in the tables are those from the best epochs. In most of the cases the models start overfitting after a certain epoch resulting in a deterioration of the accuracy after reaching a local maximum. In rare cases the test accuracy converges. Such a case can be observed in figure 8.5.

It should be clarified that the chosen configurations of models to be trained make no claim to completeness and do not claim to produce the best results, as they are not optimized individually. In all following tables the word round is abbreviated with r. and accuracy with acc.

	5 r.		10 r.		15 r.		20 r.	
	Acc.	Loss	Acc.	Loss	Acc.	Loss	Acc.	Loss
sd0x	65.12%	0.6584	81.88%	0.507	78.63%	0.5415	78.75%	0.5039
sd1x	71.25%	0.6385	81.75%	0.4899	77.38%	0.5399	80.62%	0.4893
sd2x	68.5%	0.6476	77.38%	0.5128	77.25%	0.539	80.0%	0.4965
sd3x	69.38%	0.6494	77.5%	0.5135	75.75%	0.5364	79.88%	0.4928
sd4x	67.5%	0.6436	77.5%	0.5446	76.13%	0.5418	81.12%	0.501
sd5x	72.25%	0.6212	74.5%	0.5404	78.25%	0.5286	81.0%	0.5005
sd6x	71.87%	0.6346	75.25%	0.5412	77.12%	0.5356	80.25%	0.4967
sd7x	70.25%	0.6252	73.38%	0.5453	75.63%	0.5357	81.12%	0.4918
sd8x	74.25%	0.6121	73.38%	0.5486	77.0%	0.5316	80.38%	0.5036
sd9x	74.88%	0.5826	74.25%	0.5369	76.5%	0.5408	80.25%	0.4909
sd10x	74.87%	0.5971	76.38%	0.5594	76.25%	0.548	80.25%	0.501
sd20x	73.25%	0.5831	73.25%	0.5414	76.0%	0.5387	80.87%	0.4935

Table 8.1: Best accuracies and losses for 5, 10, 15 and 20 rounds produced by the 1D convnet, prior to the changes to the simulator.

	5 r.		10 r.		15 r.		20 r.	
	Acc.	Loss	Acc.	Loss	Acc.	Loss	Acc.	Loss
sd0x	63.5%	0.6471	74.0%	0.4991	79.62%	0.4726	79.0%	0.4778
sd1x	68.12%	0.619	76.62%	0.5171	79.5%	0.4984	81.75%	0.4735
sd2x	68.25%	0.6382	80.0%	0.5243	80.38%	0.4851	84.12%	0.4573
sd3x	69.5%	0.6392	78.88%	0.5518	80.25%	0.5038	84.62%	0.4718
sd4x	67.63%	0.6415	76.38%	0.5685	79.62%	0.5106	85.0%	0.4726
sd5x	70.88%	0.6339	74.63%	0.5668	80.0%	0.51	84.88%	0.474
sd6x	70.62%	0.638	77.0%	0.5623	80.12%	0.5077	84.88%	0.4776
sd7x	69.75%	0.6543	75.38%	0.5762	79.5%	0.5176	85.0%	0.4839
sd8x	70.0%	0.6467	74.25%	0.5775	80.38%	0.5164	85.0%	0.484
sd9x	71.25%	0.6395	78.75%	0.5696	81.5%	0.5064	85.0%	0.4814
sd10x	68.75%	0.6296	77.75%	0.5674	80.12%	0.5098	85.0%	0.4802
sd20x	73.0%	0.5833	76.63%	0.5486	79.75%	0.5064	84.88%	0.4704

Table 8.2: Best accuracies and losses for 5, 10, 15 and 20 rounds produced by the 2D convnet, prior to the changes to the simulator.

	5 r.		10 r.		15 r.		20 r.	
	Acc.	Loss	Acc.	Loss	Acc.	Loss	Acc.	Loss
sd0x	63.87%	0.6698	81.88%	0.5014	78.5%	0.5401	78.62%	0.5107
sd1x	67.75%	0.6648	79.0%	0.5095	75.75%	0.5305	79.38%	0.5101
sd2x	66.0%	0.6506	79.25%	0.4845	76.0%	0.5325	78.37%	0.514
sd3x	64.38%	0.6585	77.38%	0.5037	77.62%	0.5261	79.25%	0.5113
sd4x	67.0%	0.6377	75.38%	0.5289	77.62%	0.5331	80.25%	0.5027
sd5x	67.0%	0.6452	76.75%	0.5146	78.75%	0.5226	79.0	0.5042
sd6x	70.25%	0.6427	75.25%	0.5335	78.75%	0.5208	79.38%	0.5049
sd7x	70.25%	0.6449	76.88%	0.5114	78.0%	0.5129	78.75%	0.5115
sd8x	71%	0.628	76.25%	0.523	77.25%	0.5261	79.12%	0.5156
sd9x	70%	0.6276	75.37%	0.5226	77.62%	0.5281	78.63%	0.522
sd10x	71.5%	0.6242	76.25%	0.5261	76.38%	0.5331	78.62%	0.4986
sd20x	71.88%	0.6246	72.38%	0.5617	76.25%	0.5416	79.12%	0.5

Table 8.3: Best accuracies and losses for 5, 10, 15 and 20 rounds produced by the 1D convnet, after the changes to the simulator.

	5 r.		10 r.		15 r.		20 r.	
	Acc.	Loss	Acc.	Loss	Acc.	Loss	Acc.	Loss
sd0x	63.62%	0.6461	73.87%	0.4929	80.0%	0.4557	78.5%	0.476
sd1x	67.88%	0.6446	76.12%	0.531	79.5%	0.4911	79.75%	0.4874
sd2x	69.75%	0.6463	79.62%	0.5243	80.87%	0.4905	80.0%	0.4859
sd3x	69.5%	0.6449	77.5%	0.5507	82.25%	0.4994	80.0%	0.4865
sd4x	69.0%	0.6347	72.75%	0.5684	80.12%	0.5156	80.0%	0.495
sd5x	69.38%	0.6342	73.62%	0.5634	80.5%	0.5122	80.0%	0.4917
sd6x	71.0%	0.6326	74.5%	0.5586	80.62%	0.5038	79.5%	0.4874
sd7x	70.0%	0.6328	71.62%	0.555	80.0%	0.5061	79.5%	0.4874
sd8x	69.88%	0.6287	72.75%	0.5567	80.25%	0.5057	79.62%	0.4874
sd9x	69.12%	0.6238	72.38%	0.5532	79.5%	0.5067	78.88%	0.4887
sd10x	69.5%	0.6203	73.25%	0.5536	79.5%	0.5062	79.0%	0.4904
sd20x	69.88%	0.5935	73.0%	0.5494	79.62%	0.5005	78.75%	0.4874

Table 8.4: Best accuracies and losses for 5, 10, 15 and 20 rounds produced by the 2D convnet, after the changes to the simulator.

The results show that using a certain number of simulated games for the training can often significantly improve the classification results over using only the real data. A good example is the accuracy for sd0x compared to sd4x for 20 rounds for the 2D convnet in table 8.2. Using no simulations the accuracy is 79.0% while for a ratio of 4 simulations for each real game the accuracy is 85.0%. Furthermore, it can be seen that using too many simulations compared to real games can decrease the accuracy. In table 8.2, for 10 rounds, the accuracy for sd2x is 80.0% and less when using more simulations. This behaviour highly varies depending on the configuration of the model. Sometimes the models produce the best results without any simulations, while other times using simulations strongly outperforms using only real data. It is possible that the optimal ratios of simulated to real games are not included the tested configurations. Especially the 2D convnet for 5 rounds before the changes to the simulator achieves its best accuracy with the highest ratio that was tested. Therefore, the optimal ratio could be higher than 20 simulations for each real game.

When looking at the best accuracies of 1D and 2D convnets for each number of rounds, it can be observed, that the best accuracies of the 1D convnets are lower on 15 and 20 rounds than on 10, while this is not the case for 2D convnets. This can be seen in figure 8.3. The fact that the 1D convnets do not profit from more than 10 rounds might be related to the findings from section 5.6 *Evaluation of simulation*, which indicate that the simulator performs worse in later rounds.




Figure 8.3: Best results for the different number of rounds before the changes to the simulator. Ratios of real to simulated games is additionally described.

However, if the case, this begs the questions why this does not apply to 2D convnets. There is a possibility that the worse performance in later turns has a strong impact on the statistical features fed to the 1D convnet, but once these features are used to create the synthetic images for the 2D convnet this negative impact decreases. Figure 8.3 additionally shows the ratios of simulated to real data for each of the best results. The fact that the 2D convnets for 15 rounds have the best results when using 9 simulated games for each real game while the best configurations of the 1D convnets for 15 rounds use no simulations, supports the theory from above. Nonetheless the findings are not sufficient to verify the theory. For that more research would be required. However, if found to be true, it could mean that 2D convnets using synthetic images have the potential to handle inaccurate simulations better than 1D convnets. From a different perspective, this could also indicate that, assuming the simulator does actually perform worse in later rounds and is improved in that regard, the accuracy of the 1D convnets after 10 rounds could be improved.

It was considered to perform statistical tests comparing different models and ratios to reveal significant or insignificant differences. However, the decision was made that this does not really give reliable results, because the results are not necessarily the best achievable as the models were not individually optimized. Therefore, the results of statistical tests performed on the achieved accuracies, would not justify any statements regarding whether 1D convnets or 2D convnets perform significantly better. Such statistical tests would only show significant differences or indifferences in the results, but could not be interpreted for any valuable findings, as not all models are used to their full potential. For choosing the best configurations from those tested no statistical tests are used. Instead simply those with the highest accuracies are chosen.

	5 r.	10 r.	15 r.	20 r.	Average
1D convnet - before	74.88% (sd9x)	81.88% (sd0x)	78.63% (sd0x)	81.12% (sd4x)	79.13%
1D convnet - afterwards	71.88% (sd20x)	81.88% (sd0x)	78.75% (sd5x)	80.25% (sd4x)	78.19%
2D convnet - before	73.0% (sd20x)	80.0% (sd2x)	81.5% (sd9x)	85.0% (sd4x)	79.88%
2D convnet - afterwards	71.0% (sd6x)	79.62% (sd2x)	82.25% (sd3x)	80.0% (sd2x)	78.22%

Table 8.5: Highest accuracies before and after the changes to the simulator. Furthermore, the ratios of simulated to real games and the average accuracies of the best results from each configuration are shown.

Table 8.5 shows the best result for all categories before and after changing the simulator. It can be seen that although some accuracies are slightly better after the changes to the simulator the results are on average noticeably worse than before. The accuracy from the 2D convnet for 20 rounds, for instance, is decreased from 85.0% to 80.0%. The reason for this is unknown and can be highly complex, due to the many factors involved. A theory is, that reducing the randomness of the simulator and therefore making the simulated games more similar to the real games, may have resulted in a decreased generalising ability of the models and therefore less robustness on unknown data. However, this theory can hardly be verified. The overall decreased performance after the changes leads to the decision to only use the models trained on the simulated data created before the changes to the simulator. The chosen configurations for the different numbers of rounds can be seen in table 8.6.

	5 r.	10 r.	15 r.	20 r.
1D convnet	74.88% (sd9x)	81.88% (sd0x)		
2D convnet			81.5% (sd9x)	85.0% (sd4x)

Table 8.6: The chosen configurations and their accuracies (averaged over 400 models each).

The 1D convnets trained, perform better for 5 and 10 rounds than 2D convnets, while 2D convnets perform better for 15 and 12 rounds. Depending on the number of rounds it can therefore be beneficial to use either 1D convnets or 2D convnets. The 4 best configurations for the different numbers of rounds are used in section 8.4 [Ensemble-based systems](#). Figures 8.4, 8.5, 8.6 and 8.7 show the average training and test accuracy during the course of training the models of these 4 configurations.




Figure 8.4: Average training and test accuracy of 400 1D convnets on 5 rounds and a ratio of 9 simulated games for each real game. Best test accuracy is 74.88% at epoch 812.




Figure 8.5: Average training and test accuracy of 400 1D convnets on 10 rounds using no simulated games. Best test accuracy is 81.88% at epoch 2498.




Figure 8.6: Average training and test accuracy of 400 2D convnets on 15 rounds and a ratio of 9 simulated games for each real game. Best test accuracy is 81.5% at epoch 111.




Figure 8.7: Average training and test accuracy of 400 2D convnets on 20 rounds and a ratio of 4 simulated games for each real game. Best test accuracy is 85.0% at epoch 75.

8.4 Ensemble-based systems

The accuracies shown in section 8.3 [Training results and evaluation](#) can not be used as measurement for the real life performance of the models. Since the accuracy is the mean of 20 repetitions for each split, there is no single model that produces that accuracy. A naive approach could be to simply select one of the models. However, selecting a single model for the classification comes with a problem that should not be ignored: Due to the fact that convnets are stochastic machine learning algorithms, multiple models, even if they have similar training performances, may have different generalization performances. In such cases, combining the predictions of several classifiers by letting them vote for the final classification may reduce the risk of an unfortunate selection of a poorly performing classifier. Systems that incorporate this method are called ensemble-based systems. They do not necessarily beat the performance of the best classifier in the ensemble, but they certainly reduce the overall risk of selecting a poorly performing classifier [Pol06, p. 22].

Therefore, an ensemble-based system for each of the best convnets and ratios of simulated to real data for each number of steps is created. As shown in table 8.5 the 1D convnet creates the best accuracy for 5 and 10 rounds while the 2D convnet creates the best accuracies for 15 and 20 rounds. Hence, in total 4 ensemble-based systems are created, one for each of the 4 numbers of rounds. As before, the systems are tested using leave-one-out cross validation and each of the 20 repetitions of a split is only tested with the two real games (one no obstacle and one glare) that were not included in the training of that split. However, instead of directly calculating the accuracy of each repetition of that split, the output labels are saved for all repetitions and used to perform a voting for the final label. Meaning that if for instance 15 out of 20 repetitions classified the game as a glare effect game, while the other 5 classified it as a no obstacle game, the final classification will be that of a glare effect game. This is done for the two test games specific to each split. Using the final labels created by the voting, the accuracy of that split is calculated. In the example above, if assumed that both final labels would be correct, the accuracy of that split would be 100%. This procedure is repeated for each of the 20 splits and the accuracies of the splits are averaged.

As none of the models have been saved after the training, the 400 models for each of the 4 ensemble-based systems, are retrained and saved. In total 1600 models are trained and saved for the ensemble-based systems. However, they are trained for less epochs than in section 8.3 [Training results and evaluation](#), as the models would otherwise overfit. For each of the 4 configurations the number of epochs is chosen so that the accuracy of the models in the last epoch, and therefore the accuracy of the saved model, is in the area of the maximum. The number of epochs was intentionally chosen only in a rough area of the maximum of the test accuracy and not exactly at the maximum. As mentioned in section 8.1 [Models](#), there is no validation set and therefore using the exact number of epochs as optimal for the test data would mean that the models would be optimized according to the test data and the results would be less representative of the real world performance of the systems. Table 8.7 shows amongst other things after which epoch the training of the different models was stopped and what the average accuracies were after the according epoch, with and without voting. The slight difference of the average accuracies without voting shown in the table and the best accuracies shown in table 8.6 are a result of two

factors: Firstly, the number of training epochs was only chosen roughly in the area of the maximum accuracy and secondly, due to the nature of stochastic models, even though the extend is reduced by averaging over multiple repetitions, the accuracy in a specific epoch can still vary in every training [Bro]. It should be emphasized that the average accuracies in table 8.7 are not that of the best epochs but those produced by the final model, meaning after the last training epoch.

	training epochs	avg. acc.	acc. of EBS	Overruled correct classifications	Overruled false classifications
5 r.	650	73.13%	72.5%	38	33
10 r.	2400	80.63%	82.5%	23	38
15 r.	90	80.63%	80.0%	9	4
20 r.	60	85.0%	85.0%	0	0

Table 8.7: Number of epochs the models were trained for, their average accuracy with and without voting, as well as the number of overruled correct and false classifications. EBS stands for ensemble-based system.

When comparing the average accuracies without voting with the accuracies of the ensemble-based systems, it can be seen that the accuracy of the ensemble-based systems is either better, worse or equal to the average accuracy without voting. Key factors that influence whether the ensemble-based system performs better than the average accuracies of the models are the numbers of correct and false classifications that are overruled by the majority in the voting. If more correct classifications are overruled than wrong ones, the accuracy of the ensemble-based system will be worse than the average accuracy without voting. The reason is that once a correct classification is overruled by the majority, it has no influence on the calculated accuracy any more, because the accuracy of the ensemble-based system is determined by the final labels. A good example of such a case occurred in the ensemble-based system for 5 rounds, where in one of the splits, 11 out of the 20 models voted incorrect, resulting in 9 correct predictions being ignored. Likewise, the accuracy with voting is higher than without if more false predictions are overruled than correct ones. Table 8.7 also shows how many correct and false predictions were overruled by their opposition which can directly be connected to whether the accuracy with or without voting is higher. The only voting based system that outperforms the average accuracy is that based on 1D convnets for 10 rounds, as there are more overruled false than correct predictions. Furthermore, the systems for 5 and 15 rounds perform worse than the average accuracy. Lastly, for 20 rounds, the results with and without voting are equally good.

It is important to mention that the created systems are not meant for production use, even if the fact that they are not optimized is overlooked. The training of models for production use, unlike before, should not be divided into 20 splits from which each split excludes real data for testing, but instead all of the real data should be included in the training. The validation of these systems would be done on new data collected from new participants.

Contrary to using a single model for the classification, these ensemble-based systems reduce the overall risk of making a particularly poor selection of a model [Pol06, p. 22]. Furthermore, the accuracies of these systems are more representative of their performance in real world scenarios than the accuracies in section 8.3 [Training results and evaluation](#). In the context of detecting interaction obstacles, models that use

less rounds are of higher relevance, as they make their prediction earlier. If only one of the 4 ensemble-based systems was used, the one for 10 rounds with an accuracy of 82.5% should be the preferred choice. Making the prediction after 5 rounds with an accuracy of 72.5% seems less ideal. Same goes for the two ensemble-based systems for 15 and 20 rounds. The system trained on 15 rounds is less accurate than the system for 10 rounds and waiting 20 rounds before making a prediction is too long. By then the user might have already left the game, due to a bad experience caused by the interaction obstacle. However, it is also possible to use all 4 systems. Each 5 rounds could be predicted whether an interaction obstacle is involved. This could be especially useful as the glare effect is a transient visual obstacle, unlike for instance colour blindness. Therefore, it might only be an issue for a couple turns and not for the duration of the whole game. If predicted that the sun is no issue any more, the eventual adaptation could be removed again. Continuing the adaptation even though there is no need for it any more can be harmful for the user experience and acceptance [SPIS19, p. 1].

9. Conclusion

A similarity matrix was created that successfully describes the colour differences of the cards under the influence of the glare effect. It was used in a cognitive user simulator to accurately simulate glare effect games. This simulator was also used to simulate no obstacle games. It was attempted to improve the simulator, but it could not be clearly said if that was successful. Although there are some indications that the quality of the simulations was improved, there are also others findings that suggest the opposite. One of them being that the classification results using the simulations were worse after the changes to the simulator. Additional findings suggest that the simulator performs less good in later rounds of the game. Furthermore, it was shown that incorporating certain amounts of simulated data improves the classification results significantly.

For the classification of the visual interaction obstacle being the glare effect two approaches were implemented. Firstly, a 1D convnet was utilized that classifies a sequence of statistical features. The second approach was less common: Creating synthetic images out of the naturally sequential data and use 2D convnets to classify the created images. Classifiers for four different numbers of rounds were tested. While the best result for 5 and 10 rounds are achieved by the 1D convnets, the 2D convnets outperform the 1D convnets when using 15 and 20 rounds. Findings indicate a possibility that 2D convnets are less effected by inaccurate simulations than 1D convnets.

Four ensemble-based systems, one each for 5, 10, 15 and 20 rounds were implemented, that decide the final classification by a voting of multiple models. The systems for 5 and 10 rounds utilize 1D convnets while those for 15 and 20 rounds 2D convnets. The accuracies of the four systems in the order of the number of rounds were 72.5%, 82.5%, 80.0% and 85.0%. As the topologies and the hyper parameters of the models are not individually optimized due to small number of recordings and instead best practice values were chosen it is advisable to collect data from new subjects and optimize the models before using them. Nevertheless, two different usages are suggested for such ensemble-based systems. If the requirement is to use a single system, the one trained to classify games after 10 rounds should be chosen, as the accuracy of 82.5% is high and the decision is made relatively early in the game. Otherwise an option would be

to use all four ensemble-based systems over the course of the game, depending on the number of rounds. This would mean making a prediction every 5 rounds until round 20 and adjusting the eventual adaptation accordingly. When considering that the glare effect is a transient visual obstacle, meaning that adaptation may not be necessary throughout the whole game, the second approach seems more sensible. Making a prediction each 5 rounds would bring the advantage of being able to switch the adaption on and off more often. Continuing the adaptation even though it is not needed any more can be harmful to the user experience and acceptance [SPIS19, p. 1]. Contrary to that, in the case of persistent visual obstacles, such as colour blindness, multiple tests in different stages of the game are not useful, as the adaptation, if needed once, is needed throughout the whole game.

9.1 Prospect

Regarding the simulation of user behaviour, instead of only simulating once and using a number of the best simulations, it could be simulated multiple times to take only the single best simulation from each run. This could improve the overall quality of the simulations used to train the models. Regarding the simulator, further research could be done to find out whether it actually performs worse in later turns, as indicated by findings of this work. If this is the case and the reasons are discovered, further attempts could be made to improve the simulator.

Creating synthetic images out of sequential data and using the image for an image classification is very experimental and only one way of creating such images was tested in this work. Therefore, it could be analysed how different visual representations of the sequential data influence the performance of the classifier, as it is possible that a better representation exists.

As mentioned in section 8.1 *Models*, the data was only split into training and test data without separate validation data as the amount of game logs collected is very small. This is also the main reason why the models were not optimized, as a separate validation set would be needed. The models could be optimized on the validation data and then tested on the test data, assuring that they perform well on data that they have never seen before. A new user study could be performed in order to collect the necessary data to optimize the models. Such an optimization could significantly improve the classification results.

Last but not least, it could be interesting to further investigate the creation of synthetic image and the rather unconventional usage of image recognition using 2D convnets on what was originally a sequence. In this work this approach outperformed the more conventional use of 1D convnets in multiple cases. There is a possibility that the 2D convnets were more capable in handling inaccurate simulations. Further research could potentially reveal that this is actually the case or that there are other aspects to using synthetic images and 2D convnets that make them outperform 1D convnets in certain sequence classifications. Both would be interesting findings. The first step could be to analyse whether this behaviour can still be observed after collecting additional data and optimizing the individual models.

A. Appendix

A.1 Scripts

Besides some small changes, various classes and scripts were added to existing projects. Furthermore, multiple scripts independent from those projects were written.

All scripts independent from already existing projects are located in the digital folder *ba_glare_effect_scripts* that was added to the submission. The folder contains a subfolder called *classification* as well as the following scripts:

- `similarity_matrix_generator.py`
- `similarity_matrix_mapper.py`
- `similarity_matrix_plotting.py`
- `valid_logs_collector.py`
- `feature_generation.py`
- `best_results_plotting.py`

The first three scripts were used to create the similarity matrix, add it to the existing game logs and to create figure 5.3. If the required images are provided in a resolution of 1920 x 1080, it can be used to create a similarity matrix for any type of visual obstacle. Otherwise the areas pixels are extracted from need to be adjusted. The script *valid_logs_collector.py* loads the real game logs and saves the valid ones in new files. Once the games are simulated, the script *feature_generation.py* loads all raw game logs and creates the statistical features for them, mentioned in section 5.7.1 1D convnet features. Lastly *best_results_plotting.py* was used to create figure 8.3. Some parts in *similarity_matrix_mapper.py* and *feature_generation.py* are taken from the Memory-Statistic project but most of it had to be changed or rewritten in order to be utilized for this work.

The aforementioned subfolder *classification* contains the following scripts:

- glare_effect_classification_cnn_1D.py
- glare_effect_classification_cnn_2D.py
- glare_effect_cnn_voting.py

The first two are the script for training and testing the 1D and 2D convnets, while *glare_effect_cnn_voting.py* contains the code for the ensemble-based systems.

The following additions were made to the CMM project on a new branch called *glare_effect_simulation*:

- NoObst_ConfigGenerator_dataCollection2020.java
- GlareEffect_ConfigGenerator_dataCollection2020.java
- NoObstWinStrategyTrainingDataGenerator_dataCollection2020.java
- GlareEffectWinStrategyTrainingDataGenerator_dataCollection2020.java

The first two are used to create configuration files that contain the optimized parameters for the simulation and the last two load the files and simulate the according no obstacle and glare effect games. They use functionalities from the simulator and already existing scripts served as example. The changes made in an attempt to improve the simulator can also be found on this branch.

Furthermore, the following additions were made to the Memory-Statistic project on a new branch called *Glare_Effect_Statistics*:

- memory_data_analysis_glare_effect.py
- sorting_simulated_data_quality_glare_effect.py
- similarity_evaluation_glare_effect.py
- leave_one_subject_out_glare_effect.py

The first one is simply a copy of an already existing script that contained functionalities to which minor adjustment were made so that it could be used in the context of this work. The other three use those functionalities.

sorting_simulated_data_quality_glare_effect.py was used to sort the simulations according to how close their performance is to the real games and

similarity_evaluation_glare_effect.py loads those sorted simulations as well as the real games in order to plot the two performance measurements, mentioned in section 5.6 [Evaluation of simulation](#). Furthermore, it was used to perform the various paired sample t-tests. Finally the script *leave_one_subject_out_glare_effect.py* creates the splits with the raw game logs, from which the statistical features are calculated.

Lastly, the changes to the Memory-Game project made to collect the data for the creation of the similarity matrix are uploaded on a new repository called *Memory_Game_Glare_Effect*. Uploading the changes on a new branch of the original project was not possible due to the lack of permissions.

A.2 Data

The following data is stored on the *share* server:

- screenshots and the data used for the creation of the similarity matrix
- raw game logs as well as statistical features for those games before and after the changes to the simulator
- the statically and dynamically mapped logs
- all data from the simulations
- training results before and after the changes to the simulator
- scripts used for training and testing as well as the other ones that were also created independent from existing projects
- models used in the four ensemble-based systems
- a presentation that was held regarding the two approaches for 1D and 2D convnets

Bibliography

- [ABB⁺04] John Anderson, Daniel Bothell, Michael Byrne, Scott Douglass, Christian Lebiere, and Yulin Qin. An integrated theory of the mind. *Psychological review*, 111:1036–60, 11 2004.
- [Ada20] Nikolas Adaloglou. Understanding the receptive field of deep convolutional networks. <https://theaisummer.com/>, 2020.
- [Amu] M. Amunategui. Convolutional neural networks and unconventional data - predicting the stock market using images. <http://amunategui.github.io/unconventional-convolutional-networks/>. Last accessed: 2020-09-24.
- [Bro] Jason Brownlee. Estimate the number of experiment repeats for stochastic machine learning algorithms. <https://machinelearningmastery.com/estimate-number-experiment-repeats-stochastic-machine-learning-algorithms/>. Last accessed: 2020-09-24.
- [Cho17] François Chollet. *Deep Learning with Python*. Manning, November 2017.
- [Cie] Cie2000 calculator. <http://colormine.org/delta-e-calculator/cie2000>. Last accessed: 2020-09-26.
- [Co18] François Chollet and others. Keras: The Python Deep Learning library, June 2018.
- [DHK13] L. Deng, G. Hinton, and B. Kingsbury. New types of deep neural network learning for speech recognition and related applications: an overview. In *2013 IEEE International Conference on Acoustics, Speech and Signal Processing*, pages 8599–8603, 2013.
- [Fis02] J. Fischer. Untersuchung der farbabstandsformeln des cielab farbraums auf ihre eignung, farbrauschen quantitativ und physiologisch richtig zu beschreiben, 2002.
- [Ihr19] M. Ihrig. Analyse und vergleich von benutzerverhalten und gehirnaktivität bei verschiedenen versionen eines memory-spiels zum test adaptiver hilfstellungen, 2019.
- [LLUZ17] Wenjie Luo, Yujia Li, Raquel Urtasun, and Richard S. Zemel. Understanding the effective receptive field in deep convolutional neural networks. *CoRR*, abs/1701.04128, 2017.

- [Mir18] R. Miranda. Recognising & dynamically adapting to visual and memory-based hci obstacles based on behavioural data, 2018.
- [Pas01] George Paschos. Perceptually uniform color spaces for color texture analysis: An empirical evaluation. *Image Processing, IEEE Transactions on*, 10:932 – 937, 07 2001.
- [Pol06] R. Polikar. Ensemble based systems in decision making. *IEEE Circuits and Systems Magazine*, 6(3):21–45, 2006.
- [SHK⁺14] Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. Dropout: A simple way to prevent neural networks from overfitting. *Journal of Machine Learning Research*, 15(56):1929–1958, 2014.
- [SP18a] Mazen Salous and Felix Putze. Behaviour-based working memory capacity classification using recurrent neural networks. In *ESANN 2018 – 26th European Symposium on Artificial Neural Networks, Computational Intelligence and Machine Learning*, pages 159–164, Brugge, Belgium, 2018.
- [SP18b] Mazen Salous and Felix Putze. Behaviour-based working memory capacity classification using recurrent neural networks. 04 2018.
- [SPIS19] Mazen Salous, Felix Putze, Markus Ihrig, and Tanja Schultz. Visual and memory-based hci obstacles: Behaviour-based detection and user interface adaptations analysis. 10 2019.
- [Ten20] F. Tenzer. Anzahl der smartphone-nutzer in deutschland in den jahren 2009 bis 2019. 2020.
- [To] Yosi Taguri and others. Keras conv1d: Working with 1d convolutional neural networks in keras. <https://missinglink.ai/guides/keras/keras-conv1d-working-1d-convolutional-neural-networks-keras/>. Last accessed: 2020-09-23.
- [WLW⁺16] S. Wang, W. Liu, J. Wu, L. Cao, Q. Meng, and P. J. Kennedy. Training deep neural networks on imbalanced data sets. In *2016 International Joint Conference on Neural Networks (IJCNN)*, pages 4368–4374, 2016.
- [Won15] Tzu-Tsung Wong. Performance evaluation of classification algorithms by k-fold and leave-one-out cross validation. *Pattern Recognition*, 48(9):2839 – 2846, 2015.
- [YRK⁺15] Steven R. Young, Derek C. Rose, Thomas P. Karnowski, Seung-Hwan Lim, and Robert M. Patton. Optimizing deep learning hyper-parameters through an evolutionary algorithm. In *Proceedings of the Workshop on Machine Learning in High-Performance Computing Environments, MLHPC '15*, New York, NY, USA, 2015. Association for Computing Machinery.